
The Convergence Cost of Reasoning: Why Per-Token Early Exit Underperforms

Anonymous Author(s)

Affiliation

Address

email

Abstract

Reasoning traces exhibit a 15.3pp per-token exit accuracy deficit that, combined with compound error over $\sim 3,000$ tokens, defeats all standard CALM criteria. On R1-Qwen3-8B, softmax-based criteria (response, gap) yield 0% question accuracy (Q_Acc); cosine similarity reaches at most 20% Q_Acc at negligible ($<0.1\%$) savings. Via thinking mode ablation on one architecture (Qwen3-8B), we decompose two failure mechanisms: oscillatory dynamics (distillation-specific; absent under native reasoning) and representation immaturity (reasoning-general). Tuned lens analysis reveals a +45.6pp oracle gap (in-distribution), identifying projection quality as the primary bottleneck. Despite oracle recovery, practical exit remains limited (29.4% Q_Acc vs. 84.3% full-depth), confirming compound error as the binding deployment constraint. A learned gate achieves 67–81% Q_Acc across three reasoning architecture families (Qwen, Llama, Phi) at 8B–14B scale; the 7B distilled variant and non-reasoning control show chance-level gate loss, suggesting a minimum-capacity threshold for learnable exit signals. CALM failure persists at 32B (QwQ-32B); both distilled models share a single teacher, and a native thinking-mode control confirms separability (with reversal at 1.5B; §6). Cross-benchmark replication on GSM8K confirms directional consistency (Appendix Z).

1 Introduction

Per-token early exit—one of the most promising inference optimizations for long reasoning traces—collapses to 0% question accuracy on current reasoning models. Models such as DeepSeek-R1 [DeepSeek-AI, 2025] and QwQ [Qwen Team, 2024] routinely exceed 3,000 tokens per question, and compound error across such sequences is well known [Bengio et al., 2015, Ranzato et al., 2016]: even 99.2% per-token accuracy yields 0% question-level accuracy over $\sim 3,000$ tokens. However, this observation alone does not explain *why* compound error is so severe in reasoning traces, nor does it suggest *how* to mitigate it. Reasoning traces are not homogeneous: they interleave numerical computation, logical connectives, and formatting scaffolding—token roles with fundamentally different convergence dynamics that prior work has not fully characterized. At the sequence lengths typical of translation or summarization (50–100 tokens), 99.2% per-token accuracy yields ~ 45 –67% all-correct sequence accuracy—degraded but tolerable under soft metrics; at 3,000 tokens, the same rate yields effectively 0%—a qualitative rather than quantitative shift.

Comparing 2 architectures $\times$ 2 training conditions, we decompose two failure mechanisms with architecture-dependent reversal: oscillatory dynamics in numerical tokens produce false convergence signals, while reasoning keywords exhibit representation immaturity despite appearing confident—two independent mechanisms that undermine any single exit threshold (pseudo- $R^2 \leq 0.10$, $p \approx 0$).

We term the resulting shortfall the *Reasoning Early Exit Gap*: tuned lens oracle reveals a +45.6pp gap over raw projection (§5.3). A learned gate partially closes this gap (67–81% Q_Acc across




Figure 1: Overview of our analysis framework. (a) Token roles in reasoning traces: four categories with distinct convergence properties. (b) Per-token early exit (schematic): each token can potentially exit at an intermediate layer, but reasoning tokens resist early convergence. (c) Two failure modes identified: oscillatory dynamics (primarily in distilled models) co-occur with false convergence signals, while representation immaturity (across all reasoning models) limits tokens from reaching exit thresholds at sub-final layers. Real per-token cosine-similarity trajectories from R1-Qwen3-8B on GSM8K (selected to illustrate typical failure modes; see Figure 2 right for population means).




Figure 2: **Left:** Question-level accuracy vs. compute savings for cosine-similarity exit (near-zero savings) vs. oracle (12.6% lossless, 17.6% at 95% token accuracy). **Right:** Mean cosine similarity to final layer by token role in R1-Qwen3-8B, showing non-monotonic dip in numerical tokens. Data: 500 questions, 1.49M tokens.

37 three architecture families; §5.2), though compound error remains the binding deployment constraint.
 38 These patterns extend to code reasoning (HumanEval) and 14B models.

39 Contributions.

- 40 1. **Dual failure pattern:** we identify oscillatory dynamics (distillation-specific, as confirmed on
 41 Qwen3-8B, the only architecture with a native thinking mode comparator) and representation
 42 immaturity (reasoning-general, persisting without distillation), and show that both are separable
 43 via thinking mode ablation (§4.4; one architecture).
- 44 2. **Projection bottleneck:** tuned lens reveals projection quality as the primary bottleneck (+45.6pp
 45 on GSM8K, +46.7pp on MATH-500). A learned gate achieves 67–81% Q_Acc across three
 46 three architecture families at 8B–14B scale (§5.2); compound error remains the binding deployment
 47 constraint. *Scope:* 7B–32B; CALM failure confirmed at 32B (QwQ-32B). Reverses at 1.5B (§6).

Table 1: 2×2 matched comparison. Each cell is an 8B model; cross-architecture comparisons use relative depth $d = L/L_{\text{total}}$.

	Qwen (36 layers)	Llama (32 layers)
Reasoning	DeepSeek-R1-0528-Qwen3-8B (R1-Qwen) ^a	DeepSeek-R1-Distill-Llama-8B (R1-Llama) ^b
Non-reasoning	Qwen3-8B [Qwen Team, 2025] (Qwen3) ^c	Llama-3.1-8B-Instruct (Llama) ^d

^adeepseek-ai/DeepSeek-R1-0528-Qwen3-8B ^bdeepseek-ai/DeepSeek-R1-Distill-Llama-8B
^cQwen/Qwen3-8B ^dmeta-llama/Llama-3.1-8B-Instruct

2 Related Work

Early exit for Transformers. CALM [Schuster et al., 2022] proposed confidence-based adaptive exit; we show standard CALM measures yield at most 20% Q_Acc on reasoning traces (§5). LayerSkip [Elhoushi et al., 2024] combines early exit with self-speculative decoding; subsequent work extends exit along multiple axes [Graves, 2016, Del Corro et al., 2023, Raposo et al., 2024, Fan et al., 2024, Jiang et al., 2024, Xu et al., 2025, Yang et al., 2025b, Li et al., 2024, Bae et al., 2025, Noh et al., 2026]. GateSkip [Laitenberger et al., 2026] and Wei et al. [2026] independently observe reasoning-task degradation; FlexiDepth [Luo et al., 2025] notes token-role-dependent computation demands; TIDE [Jaber and Jaber, 2026] investigates learned routers for reasoning but assumes monotonic convergence, conflicting with the oscillatory dynamics we document.¹ Song et al. [2025] document deeper-layer reliance; our contribution is a quantitative decomposition that discovers architecture-dependent reversal of two failure mechanisms.

Sequence-level reasoning efficiency. Orthogonal to per-token *layer* exit, sequence-level methods decide when to terminate chain-of-thought generation [Yang et al., 2025a, Jiang et al., 2025, Nagle et al., 2026, Sui et al., 2025, Li et al., 2025] or adapt reasoning depth to difficulty [Liu et al., 2025]; S-GRPO [Dai et al., 2025] and Step-GRPO [Chen et al., 2026a] use reinforcement learning to learn early exit policies for reasoning chains. These are complementary: they decide *how many tokens* to generate, while we analyze *at which layer* each token can exit.

Reasoning-capable LLMs. Chain-of-thought prompting [Wei et al., 2022] has been scaled through specialized training [OpenAI, 2024, DeepSeek-AI, 2025, Qwen Team, 2024], producing traces 10–100× longer than direct answers. Qian et al. [2025] identify MI-peak “thinking tokens” that drive reasoning, and Bogdan et al. [2025] show that reasoning steps contribute unequally to final answers; Lin et al. [2024] identify critical tokens via rollout sampling. Our analysis connects token roles to layer-exit feasibility.

Layer-wise representation analysis. The logit lens [nostalgebraist, 2020] projects intermediate hidden states through the LM head; the tuned lens [Belrose et al., 2024] adds per-layer affine transforms. Song et al. [2025] show reasoning relies on deeper layers; Nepal et al. [2025] find critical math-reasoning layers forged during pre-training. Layer redundancy analyses [Men et al., 2024, Yang et al., 2024] show many layers can be pruned on standard benchmarks, yet reasoning tokens resist such compression. Concurrently, DTR [Chen et al., 2026b] and DRTC [Chang, 2026] address per-token importance and critical trace segments; our work is complementary (matched-model comparison, dual failure mode decomposition).

3 Methodology

3.1 Experimental Design: 2×2 Matched Comparison

To separately characterize the associations of model architecture and reasoning-specific training, we employ a 2×2 matched comparison crossing two architectures with two training conditions at the 8B scale (Table 1; scope discussed in §6).

The reasoning models are distilled from DeepSeek-R1 (671B; R1-Qwen uses the 0528 release), producing extended traces ($\sim 3,000$ tokens per question on GSM8K). The non-reasoning models

¹TIDE and Probe and Skip [Wu et al., 2026] are concurrent work.

are instruction-tuned counterparts controlling for output quality. Our design is observational: both distilled reasoning models share a single teacher (DeepSeek-R1 671B), so the 2×2 design cannot distinguish reasoning-training effects from teacher-model effects. Our “training condition” factor captures the *net* association; the native thinking mode experiment (§4.4) confirms the deficit without distillation.

Data. We evaluate on GSM8K [Cobbe et al., 2021], MATH-500 [Hendrycks et al., 2021], HumanEval [Chen et al., 2021], and AIME 2024 (Appendix R). We profiled 100–500 questions per model; R1-Llama serves as a cross-architecture check with wider bootstrap CIs (± 1.6 pp vs. ± 0.3 pp; Appendix B). Phi-4-reasoning uses a 16K-token budget (truncation artifact ≤ 3 pp; sample details in Appendix V, N).

3.2 Token Role Taxonomy

We classify each token in a reasoning trace into four mutually exclusive categories via deterministic regex rules applied to the decoded token string:

- **Numerical:** digits, decimal points, arithmetic operators, and numeric expressions (e.g., 3.14 , $+$, $=$). These carry the computational payload.
- **Special:** punctuation, brackets, newlines, and structural markers (e.g., `\n`, `<think>`, `,`). These provide formatting scaffolding.
- **Reasoning keyword:** logical connectives and reasoning markers (e.g., “therefore”, “because”, “hence”, “so”). These signal reasoning transitions.
- **Other:** all remaining tokens, predominantly natural language content.

Token role is a statistically significant but modest correlate of exit outcomes (pseudo- $R^2 \leq 0.10$, $p \approx 0$; Appendix Y); critically, effect directions reverse across architectures (numerical OR=6.76 in R1-Qwen vs. 0.20 in Phi-4), so no universal per-type exit policy can serve both—motivating the learned gate (§5.2). The taxonomy is a decomposition tool for exposing architecture-dependent failure modes (§4), not a predictive framework.

3.3 Metrics

Non-monotonic convergence rate. For each token t , let $s_\ell = \cos(\mathbf{h}_\ell^t, \mathbf{h}_{\text{final}}^t)$ denote the cosine similarity between the hidden state at layer ℓ and the final layer. We maintain a running maximum $M_\ell = \max_{j \leq \ell} s_j$ and flag token t as *non-monotonic* if $\exists \ell : M_{\ell-1} - s_\ell > \theta$ (default $\theta = 0.03$; sensitivity analysis in Appendix A)—i.e., the similarity at any layer falls below its historical peak by more than θ . Such tokens violate the monotonic convergence assumption underlying all confidence-based exit methods. All confidence intervals are computed via question-level bootstrap resampling (1,000 iterations).

Exit accuracy. At each checkpoint layer L , we project the intermediate hidden state through the final LM head: $\hat{y}_L = \arg \max(\mathbf{W}_{\text{lm}} \cdot \mathbf{h}_L)$. *Token-level exit accuracy* is the fraction of tokens where $\hat{y}_L = \hat{y}_{\text{final}}$. *Question-level exit accuracy* (hereafter Q_Acc) requires *all* tokens in a question’s reasoning trace to match. A segment-level analysis (Appendix S) confirms this is not merely a strict criterion: even evaluating only the final boxed answer, question accuracy drops from 84.3% (full depth) to $< 2\%$ at L32; compound error across multi-token sequences is the primary bottleneck.

Proxy validation. We compute Spearman ρ between per-token convergence speed (earliest layer achieving $\cos \geq 0.95$) and binary exit correctness at each layer. An effective proxy requires strong positive ρ . With 96 tests (4 models $\times$ 6 layers $\times$ 4 roles), we apply Benjamini–Hochberg FDR correction at $\alpha = 0.05$; all reported correlations remain significant at $q < 0.01$.

Relative depth. Cross-architecture comparisons use normalized depth $d = L/L_{\text{total}}$ (e.g., L28 = $d_{0.88}$ for 32-layer, $d_{0.78}$ for 36-layer models).

Table 2: Non-monotonicity rates ($\theta = 0.03$) across the 2×2 design, Phi-4-reasoning 14B cross-architecture validation ($n=180$, 2.8M tokens), and native thinking mode ablation. Token-level rates; ANOVA uses question-level means (4×100 balanced, excludes thinking mode and Phi-4).

Model	Training	Overall	Numerical	Special	Reas. KW	Other
R1-Qwen3-8B	reasoning	20.8% *	42.9%	27.4%	4.2%	11.8%
R1-Llama	reasoning	19.3%	11.7%	34.1%	11.7%	11.5%
Qwen3-8B	non-reas.	11.6%	15.3%	18.9%	7.8%	8.5%
Llama-3.1-8B	non-reas.	8.5%	1.6%	8.2%	11.8%	9.6%
<i>Cross-architecture validation (14B, 40 layers)</i>						
Phi-4-reasoning	reasoning	25.5%	5.6%	43.6%	3.0%	22.3%
Phi-4 (base)	non-reas.	11.6%	8.1%	7.9%	1.2%	14.0%
Qwen3-8B (thinking)	native reas.	9.1%	16.0%	8.3%	11.5%	7.3%

*Pooled rate across GSM8K ($n=750$) and MATH-500 ($n=364$).

4 Findings

4.1 Non-Monotonic Convergence

A token’s cosine similarity trajectory $\cos(\mathbf{h}_L, \mathbf{h}_{\text{final}})$ is *non-monotonic* if it dips by $\geq \theta$ before recovering—indicating that the representation moves *away* from its final state at intermediate layers. Non-monotonicity is one observable symptom associated with reasoning training; §4.3 shows it is associated with—but does not fully account for—the exit accuracy deficit, as a second, distinct pattern is also observed.

At the 8B–14B scale, reasoning training is consistently associated with higher non-monotonicity (Table 2), but with architecture-dependent magnitude: Llama more than doubles ($8.5\% \rightarrow 19.3\%$), while Qwen increases modestly ($11.6\% \rightarrow 20.8\%$). Within our design sharing a common teacher, training condition is the stronger correlate of non-monotonicity (Appendix C); this association captures the net effect of reasoning training, distillation, and teacher-specific factors—it is not a causal estimate.

Reasoning-trained models exhibit $\sim 2\times$ the non-monotonicity of non-reasoning counterparts (Figure 3, left; stable across three disjoint subsets, max $\Delta < 1\text{pp}$; Appendix A). The pattern holds across all $\theta \in \{0.01-0.10\}$ (Appendix A). Non-monotonicity is sharply stratified by token role but the dominant role varies: in R1-Qwen, **numerical tokens reach 42.9%** vs. reasoning keywords 4.2% (**10.2 $\times$ gap**); in Phi-4-reasoning, **special tokens lead at 43.6%**. Per-type rankings are model-specific, precluding universal exit policies; robust after controlling for chain length (Appendix E, W, A). Effect directions **reverse across architectures**: numerical OR=6.76 in R1-Qwen but 0.20 in Phi-4, confirming training-paradigm-dependent mechanisms (logistic regression details in Appendix Y).

Cross-architecture and scale validation. Phi-4-reasoning (14B) confirms the pattern: non-monotonicity more than doubles versus its base (25.5% vs. 11.6%; Table 2), though the dominant token role differs (special vs. numerical). At 1.5B scale, this pattern reverses (Appendix V): both the non-monotonicity gap and exit accuracy deficit invert, constraining generalization to 8B–14B.

4.2 Exit Accuracy Deficit

Reasoning training is associated with lower token-level exit accuracy across both architectures studied at the 8B–14B scale.

Reasoning-trained models consistently show lower exit accuracy (Table 3): R1-Qwen (32.4% [32.1, 32.7]) trails Qwen3-8B by 15.3pp, R1-Llama trails Llama (74.3% [72.8, 76.0]) by 6.4pp (non-overlapping 95% bootstrap CIs). The magnitude is architecture-dependent—Qwen’s larger deficit is consistent with L28 at 78% vs. 88% relative depth. Token role significantly predicts exit correctness at every layer in R1-Qwen3-8B ($p \approx 0$), with the effect attenuating at deeper layers (Appendix Y).

Table 3: Token-level exit accuracy (question-level mean) at L28 across the 2×2 design. L28 = 78% depth for Qwen, 88% for Llama; Δ controls for depth within architecture.

	Qwen (L28 = $d_{0.78}$)	Llama (L28 = $d_{0.88}$)
Reasoning	32.4%	67.9%
Non-reasoning	47.7%	74.3%
Δ (reasoning – non-reasoning)	–15.3pp	–6.4pp

14B validation and per-type patterns. Phi-4-reasoning (14B) confirms the deficit (–8.3pp at L8; Appendix V), with a phase transition at ~ 40 –50% depth mirroring the 8B pattern. Reasoning training compresses per-type exit accuracy from a 26.3pp range to a 5.2pp band; reasoning keywords show the largest drop (–28.8pp) yet the lowest non-monotonic rate (4.2%)—smooth but insufficiently deep convergence, the second failure mode of §4.3 (details in Appendix O, O).

Tuned lens (per-layer affine transforms; Belrose et al. 2024; 400 train / 500 test questions) confirms the deficit is genuine, not an artifact of raw projection. Tuned lens raises absolute exit accuracy by +35–54pp (e.g., L24: 11.7%→65.7%), identifying raw projection quality as the primary bottleneck (see §5.3 for distribution analysis). Despite this improvement, the reasoning–non-reasoning gap *persists* at every layer (4–6pp at L24–L32; Appendix T)—the deficit is a property of reasoning representations themselves.

4.3 Proxy Failure and Dual Failure Pattern

The standard cosine-similarity proxy fails to predict exit correctness in reasoning models: for R1-Qwen, $|\rho| < 0.15$ at every layer (Figure 3, right), confirming that cosine similarity is uninformative about exit correctness. R1-Llama’s overall $\rho = +0.32$ at L24 masks a per-type inversion—reasoning keywords show $\rho = -0.51$. An entropy-based alternative is complementarily blind: reasoning keywords have low entropy yet the worst exit accuracy.

Two independent failure patterns. Two empirically distinct failure patterns characterize the exit accuracy deficit, each affecting different token populations: **(1) Oscillatory dynamics** (distilled models): numerical and special tokens exhibit non-monotonic trajectories—appearing converged then diverging—more pronounced in distillation-trained models (§4.1) but largely absent in native reasoning (§4.4); **(2) Representation immaturity** (all reasoning models): reasoning keyword tokens converge monotonically but fail to reach sufficient accuracy at sub-final layers (L32: 55.2%), with low entropy falsely suggesting confidence. No single threshold resolves both: lowering τ triggers false exits from oscillating tokens, while raising τ blocks insufficiently-converged tokens indefinitely (Appendix D). The independence of these failure patterns is confirmed by the native reasoning experiment (§4.4).

4.4 Native Thinking Mode Ablation

A native thinking mode ablation (Qwen3-8B only; Llama-3.1-8B lacks a native thinking mode, $n=499$) decomposes the two failure modes. Oscillatory dynamics are distillation-specific: thinking NM 9.1% is *below* non-reasoning 11.6%, compared to R1’s 20.8% (Table 2)—the elevated oscillation is absent without distillation. Representation immaturity persists without distillation (L28 exit accuracy 34.7% $\approx$ R1’s 32.4%, both $\ll$ non-reasoning 47.7%; Appendix U), confirming it as reasoning-general. The ANOVA η^2 (§4.1) cannot disambiguate training from teacher identity; this ablation provides a partial control on one architecture (Qwen only; cross-architecture generalization requires future work).

5 End-to-End Early Exit Analysis

We now test progressively stronger exit strategies on R1-Qwen3-8B (500 GSM8K questions), from standard confidence methods (§5.1) to a learned gate (§5.2), with tuned lens as diagnostic upper bound (§5.3).




Figure 3: **Left:** Non-monotonic convergence rates by token role across the 2×2 design (R1-Qwen3-8B, R1-Llama-8B, Qwen3-8B, Llama-3.1-8B; 500 questions each except Qwen3 $n=499$). **Right:** Cosine-similarity proxy Spearman ρ heatmap. R1-Llama’s moderate $\rho=0.32$ masks a per-type inversion ($\rho=-0.51$ for reasoning keywords).

Table 4: Adaptive exit under raw CALM criteria and tuned lens. Softmax criteria: $n=50$; cosine similarity and tuned lens: R1-Qwen $n=500$, R1-Llama $n=500$.

Projection	τ	R1-Qwen (500q)			R1-Llama (500q)		
		Tok Acc	Q Acc	Savings	Tok Acc	Q Acc	Savings
Raw (softmax resp.) [†]	0.99	86.9%	0%	8.1%	—	—	—
Raw (softmax gap) [†]	0.99	88.9%	0%	7.2%	—	—	—
Raw (cos-sim)**	0.95	94.7%	0%	0.39%	—	—	—
Raw (cos-sim)**	0.97	99.2%	0%	0.054%	—	—	—
Raw (per-type)	mixed [†]	95%	0%	0.34%	—	—	—
Tuned lens	0.95	97.0%	7.2%	18.6%	98.6%	20.2%	15.4%
Tuned lens	0.99	99.0%	29.4%	11.4%	99.5%	52.0%	9.3%

[†]“mixed”: per-type optimized thresholds. R1-Qwen: 36 layers, 9 exit points; R1-Llama: 32 layers, 6 exit points.

[‡]CALM criteria ($n=50$); Q Acc 0% at all $\tau \in \{0.5, 0.7, 0.8, 0.9, 0.95, 0.99, 0.999, 0.9999\}$.

**Exit-to-final-layer cosine similarity ($n=500$), adapted from CALM [Schuster et al., 2022]; Q Acc 0% at $\tau \geq 0.95$. At $\tau=0.9$ ($n=50$): 20% Q Acc, 0.03% savings.

R1-Llama “—”: raw CALM criteria evaluated on R1-Qwen only ($n=50$); softmax-based 0% Q Acc on R1-Qwen ($n=50$); not evaluated on R1-Llama (R1-Llama: 52.0% Q Acc vs. 0% raw).

5.1 Compound Error and Reasoning-Specific Amplification

Under raw logit lens, question-level accuracy is 0% at every layer for all four models (Table 48), regardless of exit strategy. Even at 89% depth with 73% token accuracy, at least one token per question is incorrect; evaluating only the final boxed answer, accuracy still drops from 84.3% to <2% at L32 (Appendix S).


Two levels of failure compound: *compound error* over $\sim 3,000$ tokens (even 99.2% per-token accuracy yields near-zero Q_Acc; Table 45) and a reasoning-specific *per-token deficit* (-15.3 pp; Table 3). Even boxed-answer-only evaluation collapses to 1.6% (Appendix S).

CALM-style adaptive exit [Schuster et al., 2022] fares no better (Table 4): softmax criteria yield 0% Q_Acc across all thresholds; cosine similarity at $\tau=0.97$ achieves 99.2% token accuracy yet 0% Q_Acc. Per-type optimized thresholds yield only +0.06% over uniform. Logistic probes achieve AUC 0.78–0.87 (Table 49) but AUC *decreases* with depth—deployed as exit gates, they yield 23% Q_Acc at 0.054% savings (Appendix X).

Convergence hardens within traces. Never-converging tokens increase as reasoning progresses; answer-segment tokens exhibit the highest non-monotonic rate despite converging earliest (Appendix W, S).

5.2 Projection Bottleneck and Proof-of-Concept Recovery

A +45.6pp gap between raw (12.6%) and tuned (58.2%) oracle savings (Figure 4a) localizes the bottleneck to projection quality. As a proof-of-concept, we train a learned MLP gate on tuned-lens



(a) Raw vs. tuned oracle gap (+45.6pp). (b) Oracle savings vs. token acc. Tuned lens: Figure 4: End-to-end exit analysis on R1-Qwen3-8B (300q) → 58.2% Q_Acc at 20.2% savings (5-fold stratified CV; Appendix P); net savings after probe overhead remain ~19.2% (Appendix T.1). Margin and top- k mass drive exit decisions ($|\text{coef}| > 0.5$), while type identity contributes minimally ($|\text{coef}| < 0.02$; per-type gates: V2c 64%, 18.1% savings). Compound error remains the binding constraint: 68.2% vs. $\approx 91\%$ near-full-depth Q_Acc.

confidence features² to test whether this gap can be partially closed. The gate achieves **68.2% Q_Acc at 20.2% savings** (5-fold stratified CV; Appendix P); net savings after probe overhead remain ~19.2% (Appendix T.1). Margin and top- k mass drive exit decisions ($|\text{coef}| > 0.5$), while type identity contributes minimally ($|\text{coef}| < 0.02$; per-type gates: V2c 64%, 18.1% savings). Compound error remains the binding constraint: 68.2% vs. $\approx 91\%$ near-full-depth Q_Acc.

Gap-closing progression. Four tiers emerge: (i) raw CALM: 0% Q_Acc; (ii) tuned-lens CALM ($\tau=0.99$): 29.4% Q_Acc, 11.4% savings; (iii) V2a gate ($\theta=0.99$): 68.2% Q_Acc, 20.2% savings; (iv) oracle ceiling: 58.2% savings. Each tier closes a distinct component: projection quality (i→ii), learned boundaries (ii→iii), residual compound error (iii→iv).

Controls and scale. A matched non-reasoning control (Qwen2.5-7B-Instruct, $n=500$) confirms reasoning-specificity: gate val_loss plateaus at ≈ 0.50 (chance), yielding only 5.7% savings vs. 20.2% (3.5 $\times$). QwQ-32B ($n=500$) extends CALM to 32B: $\tau=0.95$ exits zero tokens; $\tau=0.90$ collapses to 2.4% Q_Acc—all-or-nothing failure persists at scale. R1-Distill-Qwen-7B gate also fails (val_loss ≈ 0.50 , 4.5% savings; Appendix V).

Cross-architecture V2a transfer. To test whether the V2a gating signal generalizes beyond the training architecture, we evaluate the same gate on R1-Llama-8B and Phi-4-reasoning-14B (both $n=500$, GSM8K, 5-fold CV). At $\theta=0.99$: R1-Llama achieves $67.4\% \pm 5.5\%$ Q_Acc at $9.08\% \pm 0.59\%$ savings; Phi-4 achieves $80.8\% \pm 3.3\%$ Q_Acc at $6.56\% \pm 0.62\%$ savings. Gate accuracy transfers across all three reasoning architecture families (67–81% Q_Acc at $\theta=0.99$), while savings vary with architecture (6.6–20.2%). Validation loss separates the models: all three $\geq 8B$ reasoning models yield val_loss well below 0.50 (R1-Qwen $\ll 0.50$; Phi-4 ≈ 0.41), whereas R1-Distill-Qwen-7B (≈ 0.506 ; distilled from the same R1 teacher as R1-Llama) and the non-reasoning control (Qwen2.5-7B: ≈ 0.50) remain at chance level. Since R1-Llama (8B, distilled) succeeds while R1-Distill-Qwen (7B, same teacher) fails, the discriminant likely involves model capacity rather than training methodology alone. Per-architecture calibration would recover efficiency (Appendix V).

5.3 Tuned Lens as Diagnostic Upper Bound

Tuned lens achieves non-zero question accuracy: a token exits at the shallowest layer where tuned-lens confidence exceeds τ (per-layer affine transforms from §4.2; 400/500 train/test; Table 4).

²Per-layer affine transforms fitted on all 500 MATH-500 questions prior to the gate train/test split (2 parameters per layer).

Question accuracy reaches **29.4% at $\tau = 0.99$** (7.2% at $\tau \geq 0.95$), demonstrating that improved projection alone partially overcomes the per-token deficit. The oracle curve (Figure 4b) confirms projection quality as the bottleneck: tuned lens raises lossless savings from 12.6% to 58.2% (+45.6pp; generalizes to HumanEval, Appendix Q).

Cross-difficulty transfer. On MATH-500, tuned lens oracle drops to 35.4% (−22.8pp) while raw oracle remains stable, confirming the gap resides in the learned projection. **In-distribution retraining recovers it:** MATH-500 tuned oracle reaches 57.1% (+46.7pp), nearly matching GSM8K (58.2%)—distribution mismatch, not intrinsic difficulty. Despite oracle recovery, CALM Q_Acc remains at 18%, confirming compound error as the binding constraint (Appendix R; GSM8K replication in Appendix Z).

Cross-architecture and scale validation. R1-Llama confirms at $\tau=0.99$: 52.0% Q_Acc (9.3% savings; Table 4); oracle gap +36.9pp. Phi-4 (14B) shows +49.3pp oracle gap (27.3%→76.6%; Table 39), though tuned-lens CALM achieves only 2.8% Q_Acc (Appendix T). The projection bottleneck generalizes across all three families (+36.9–49.3pp); QwQ-32B and R1-Distill-Qwen-7B in §5.2. The V2a gate confirms cross-architecture transfer (67–81% Q_Acc at $\theta=0.99$; §5.2).

6 Conclusion

Two separable failure patterns—oscillatory dynamics (distillation-specific) and representation immaturity (reasoning-general)—are confirmed via thinking-mode ablation and at 14B scale; CALM failure persists at 32B. Tuned lens narrows the Reasoning Early Exit Gap by +36.9–49.3pp across three architecture families; a learned gate achieves 67–81% Q_Acc at 6.6–20.2% savings (R1-Qwen, R1-Llama, Phi-4; 8B–14B), with the 7B distilled variant at chance-level val_loss and compound error as the binding constraint.

Implications. The taxonomy’s value is diagnostic rather than predictive—type identity contributes minimally to exit decisions ($|\text{coef}| < 0.02$; Appendix P). Segment-level methods [Yang et al., 2025a] and hierarchical verification are promising paths to close the oracle-to-deployment gap.

Limitations. Both distilled models share one teacher (DeepSeek-R1); cross-architecture V2a transfer tested on two non-Qwen models only (R1-Llama 8B, Phi-4 14B). Scale: 7B–32B with **reversal at 1.5B** (Appendix V); QwQ-32B tested under CALM diagnostics only (no V2a gate or non-monotonicity profiling). Domain: math/code only. The projection bottleneck may not dominate at larger scales or cross-distribution.

References

- Sangmin Bae, Yujin Kim, Reza Bayat, Sungnyun Kim, Jiyoun Ha, Tal Schuster, Adam Fisch, Hrayr Harutyunyan, Ziwei Ji, Aaron Courville, and Se-Young Yun. Mixture-of-recursions: Learning dynamic recursive depths for adaptive token-level computation. In *Advances in Neural Information Processing Systems*, 2025.
- Nora Belrose, Zach Furman, Logan Smith, Danny Halawi, Igor Ostrovsky, Lev McKinney, Stella Biderman, and Jacob Steinhardt. Eliciting latent predictions from transformers with the tuned lens. In *International Conference on Learning Representations*, 2024.
- Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. Scheduled sampling for sequence prediction with recurrent neural networks. In *Advances in Neural Information Processing Systems*, 2015.
- Paul C. Bogdan, Uzay Macar, Neel Nanda, and Arthur Conmy. Thought anchors: Which LLM reasoning steps matter? *arXiv preprint arXiv:2506.19143*, 2025. NeurIPS 2025 Mechanistic Interpretability Workshop (Spotlight).
- Waldemar Chang. Directional reasoning trajectory change (DRTC): Identifying critical trace segments in reasoning models. *arXiv preprint arXiv:2602.15332*, 2026.
- Benteng Chen, Weida Wang, Shufei Zhang, Mingbao Lin, and Min Zhang. Step-GRPO: Internalizing dynamic early exit for efficient reasoning. *arXiv preprint arXiv:2604.16890*, 2026a.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Pondé de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Wei-Lin Chen, Liqian Peng, Tian Tan, Chao Zhao, Blake JianHang Chen, Ziqian Lin, Alec Go, and Yu Meng. Think deep, not just long: Measuring LLM reasoning effort via deep-thinking tokens. *arXiv preprint arXiv:2602.13517*, 2026b.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Muzhi Dai, Chenxu Yang, and Qingyi Si. S-GRPO: Early exit via reinforcement learning in reasoning models. *arXiv preprint arXiv:2505.07686*, 2025.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Luciano Del Corro, Allie Del Giorno, Sahaj Agarwal, Bin Yu, Ahmed Hassan Awadallah, and Subhadrata Mukherjee. SkipDecode: Autoregressive skip decoding with batching and caching for efficient LLM inference. *arXiv preprint arXiv:2307.02628*, 2023.
- Mostafa Elhoushi, Akshat Shrivastava, Diana Liskovich, Basil Hosmer, Bram Wasti, Liangzhen Lai, Anas Mahmoud, Bilge Acun, Saurabh Agarwal, Ahmed Roman, et al. LayerSkip: Enabling early exit inference and self-speculative decoding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.
- Siqi Fan, Xin Jiang, Xiang Li, Xuying Meng, Peng Han, Shuo Shang, Aixin Sun, Yequan Wang, and Zhongyuan Wang. Not all layers of LLMs are necessary during inference. *arXiv preprint arXiv:2403.02181*, 2024.
- Alex Graves. Adaptive computation time for recurrent neural networks. *arXiv preprint arXiv:1603.08983*, 2016.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. MATH: Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems*, 2021.

329 Jaber Jaber and Osama Jaber. TIDE: Token-informed depth execution for per-token early exit in
330 LLM inference. *arXiv preprint arXiv:2603.21365*, 2026. Concurrent work.

331 Guochao Jiang, Guofeng Quan, Zepeng Ding, Ziqin Luo, Dixuan Wang, and Zheng Hu. FlashThink:
332 An early exit method for efficient reasoning. *arXiv preprint arXiv:2505.13949*, 2025.

333 Yikun Jiang, Huanyu Wang, Lei Xie, Hanbin Zhao, Zhang Chao, Hui Qian, and John C. S. Lui.
334 D-LLM: A token adaptive computing resource allocation strategy for large language models. In
335 *Advances in Neural Information Processing Systems*, 2024.

336 Filipe Laitenberger, Dawid Kopiczko, Cees G.M. Snoek, and Yuki M. Asano. What layers when:
337 Learning to skip compute in LLMs with residual gates. In *International Conference on Learning*
338 *Representations*, 2026.

339 Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE: Speculative sampling requires
340 rethinking feature uncertainty. In *International Conference on Machine Learning*, 2024.

341 Zheng Li, Qingxiu Dong, Jingyuan Ma, Di Zhang, and Zhifang Sui. SelfBudgeter: Adaptive token
342 allocation for efficient LLM reasoning. *arXiv preprint arXiv:2505.11274*, 2025.

343 Zicheng Lin et al. Critical tokens matter: Token-level contrastive estimation enhances LLM’s
344 reasoning capability. *arXiv preprint arXiv:2411.19943*, 2024.

345 Xiang Liu, Xuming Hu, Xiaowen Chu, and Eunsol Choi. DiffAdapt: Difficulty-adaptive reasoning
346 for token-efficient LLM inference. *arXiv preprint arXiv:2510.19669*, 2025.

347 Xuan Luo, Weizhi Wang, and Xifeng Yan. Adaptive layer-skipping in pre-trained LLMs. In
348 *Conference on Language Modeling (COLM)*, 2025.

349 Xin Men, Mingyu Xu, Qingyu Zhang, Bingning Wang, Hongyu Lin, Yaojie Lu, Xianpei Han, and
350 Le Sun. ShortGPT: Layers in large language models are more redundant than you expect. *arXiv*
351 *preprint arXiv:2403.03853*, 2024.

352 Alliot Nagle, Jakhongir Saydaliev, Dhia Garbaya, Michael Gastpar, Ashok Vardhan Makkuva, and
353 Hyeji Kim. TERMINATOR: Learning optimal exit points for early stopping in chain-of-thought
354 reasoning. *arXiv preprint arXiv:2603.12529*, 2026.

355 Aadim Nepal, Safal Shrestha, Anubhav Shrestha, Minwu Kim, Jalal Naghiyev, Ravid Shwartz-Ziv,
356 and Keith Ross. Layer importance for mathematical reasoning is forged in pre-training and
357 invariant after post-training. In *NeurIPS 2025 Workshop on MATH-AI*, 2025.

358 Kanghyun Noh, Jinheon Choi, and Yulhwa Kim. QTALe: Quantization-robust token-adaptive layer
359 execution for LLMs. *arXiv preprint arXiv:2602.10431*, 2026.

360 nostalgebraist. Interpreting GPT: The logit lens. Blog post, <https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens>, 2020.

361 OpenAI. OpenAI o1 system card. *arXiv preprint arXiv:2412.16720*, 2024.

362 Chen Qian, Dongrui Liu, Haochen Wen, Zhen Bai, Yong Liu, and Jing Shao. Demystifying reasoning
363 dynamics with mutual information: Thinking tokens are information peaks in LLM reasoning. In
364 *Advances in Neural Information Processing Systems*, 2025. *arXiv:2506.02867*.

365 Qwen Team. QwQ: Reflect deeply on the boundaries of the unknown. Blog post, <https://qwenlm.github.io/blog/qwq-32b-preview/>, 2024.

366 Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

367 Marc’Aurelio Ranzato, Sumit Chopra, Michael Auli, and Wojciech Zaremba. Sequence level training
368 with recurrent neural networks. In *International Conference on Learning Representations*, 2016.

369 David Raposo, Sam Ritter, Blake Richards, Timothy Lillicrap, Peter Conway Humphreys, and Adam
370 Santoro. Mixture-of-depths: Dynamically allocating compute in transformer-based language
371 models. *arXiv preprint arXiv:2404.02258*, 2024.

374 Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q Tran, Yi Tay, and
375 Donald Metzler. Confident adaptive language modeling. In *Advances in Neural Information*
376 *Processing Systems*, 2022.

377 Xinyuan Song, Keyu Wang, Pengxiang Li, Lu Yin, and Shiwei Liu. Demystifying the roles of LLM
378 layers in retrieval, knowledge, and reasoning. *arXiv preprint arXiv:2510.02091*, 2025.

379 Yang Sui et al. Stop overthinking: A survey on efficient reasoning for large language models. *arXiv*
380 *preprint arXiv:2503.16419*, 2025.

381 Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc
382 Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In
383 *Advances in Neural Information Processing Systems*, 2022.

384 Rui Wei, Rui Du, Hanfei Yu, Devesh Tiwari, Jian Li, Zhaozhuo Xu, and Hao Wang. The diminishing
385 returns of early-exit decoding in modern LLMs. *arXiv preprint arXiv:2603.23701*, 2026.

386 Zimeng Wu, Donghao Wang, Chaozhe Jin, Jiaxin Chen, and Yunhong Wang. Probe and
387 skip: Self-predictive token skipping for efficient long-context LLM inference. *arXiv preprint*
388 *arXiv:2601.13155*, 2026.

389 Jiaming Xu, Jiayi Pan, Yongkang Zhou, Siming Chen, Jinhao Li, Yaoxiu Lian, Junyi Wu, and Guohao
390 Dai. SpecEE: Accelerating large language model inference with speculative early exiting. In
391 *International Symposium on Computer Architecture (ISCA)*, 2025. arXiv:2504.08850.

392 Chenxu Yang, Qingyi Si, Yongjie Duan, Zheliang Zhu, Chenyu Zhu, Qiaowei Li, Minghui Chen,
393 Zheng Lin, and Weiping Wang. Dynamic early exit in reasoning models. *arXiv preprint*
394 *arXiv:2504.15895*, 2025a.

395 Rubing Yang et al. SpecExit: Accelerating large reasoning model via speculative exit. *arXiv preprint*
396 *arXiv:2509.24248*, 2025b.

397 Yifei Yang, Zouying Huang, and Ningke Liu. LaCo: Large language model pruning via layer collapse.
398 In *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024.

399 A Taxonomy Robustness Validation

400 We assess the robustness of the four-category token role taxonomy (§3.2) along three axes: keyword
401 set sensitivity, boundary ambiguity, and unsupervised clustering structure.

402 **Reasoning keyword list.** The 28 reasoning keywords are matched via case-insensitive regex on
403 decoded token strings: *therefore, because, so, implies, let, assume, suppose, hence, thus, since, given,*
404 *if, then, means, conclude, verify, check, consider, note, recall, know, think, wait, actually, but, however,*
405 *instead, alternatively.* Tokens matching any keyword are classified as reasoning keywords unless
406 they also match the numerical or special category (which take priority).

407 **Keyword perturbation.** To test sensitivity to this keyword set, we randomly replace 5 of 28
408 keywords per trial (10 seeds). Table 5 shows that numerical, special, and other rates are invariant
409 (by construction), while the reasoning keyword rate varies from 0.2% to 5.3% (mean 3.6%, baseline
410 4.2%). The ~ 5 pp range reflects boundary words (e.g., “but”, “let”) that straddle reasoning and
411 general usage. Crucially, the *ranking* of non-monotonicity rates across types is preserved under all
412 perturbations: numerical $\gg$ special $>$ other $>$ reasoning keyword.

413 **Boundary ambiguity.** We analyze tokens at category boundaries—those whose classification could
414 change under reasonable alternative rules. Reclassifying all ambiguous reasoning keywords as “other”
415 shifts the reasoning keyword rate from 4.19% to 0.16% (-4.03 pp) while changing the other rate by
416 only -0.41 pp. Reclassifying special-boundary tokens as numerical shifts both rates by <1.5 pp. In
417 all worst-case scenarios, the cross-type non-monotonicity ranking is preserved.

Table 5: Non-monotonicity rate (%) under keyword set perturbation (10 seeds, 5 swaps each). Numerical and special rates are unaffected; reasoning keyword rate ranges from 0.2% to 5.3%.

	Numerical	Special	Reasoning KW	Other
Baseline	42.93	27.35	4.19	11.75
Mean (10 seeds)	42.93	27.35	3.64	11.70
Min	42.93	27.35	0.21	11.52
Max	42.93	27.35	5.32	11.77

Table 6: Enrichment ratios (observed/expected) for k -means clusters ($k=3$) over convergence trajectories. Values >1 indicate over-representation. Non-mono rate is within-cluster non-monotonicity.

Cluster (non-mono rate)	Numerical	Special	Reasoning KW	Other
Low-conv (24.2%)	0.98	1.14	1.20	0.93
Mid-conv (21.1%)	1.03	0.81	0.50	1.12
Early-conv (9.7%)	0.97	1.25	2.19	0.80

Unsupervised clustering structure. We apply k -means ($k=3$) to per-token convergence trajectories (cosine similarity curves across layers) without using the regex taxonomy labels. Table 6 reports the enrichment ratio (observed/expected) of each regex category in each cluster. The low-convergence cluster (highest non-monotonicity, 24.2%) is enriched for special tokens ($1.14\times$) and reasoning keywords ($1.20\times$), while the mid-convergence cluster (21.1%) is enriched for other tokens ($1.12\times$) and numerical tokens ($1.03\times$). Reasoning keywords are strongly enriched ($2.19\times$) in the early-convergence cluster (lowest non-monotonicity, 9.7%), consistent with their smooth but immature convergence pattern (representation immaturity) identified in §4.3. Note that overall NMI (Normalized Mutual Information; 0.014) and ARI (Adjusted Rand Index; 0.032) between clusters and token roles are low, indicating that token role and convergence dynamics are largely orthogonal dimensions—the enrichment patterns above reflect relative tendencies, not categorical alignment. Formal predictive validation is provided by logistic regression (Appendix Y).

Question subset stability. To rule out single-sample artifacts, we run the R1-Qwen H1 analysis on three disjoint 100-question GSM8K subsets (questions 0–99, 100–199, 200–299; greedy decoding, no randomness beyond question selection). Table 7 reports non-monotonicity rates at three thresholds.

The per-type non-monotonicity ranking is identical across all three subsets: numerical $\gg$ special $>$ other $>$ reasoning keyword (at $\theta=0.03$, numerical ranges 32–36%, reasoning keyword 2.8–3.1%). Layer-wise dip concentration is also stable: layers 16–20 account for 66–71% of all cosine similarity dips across subsets.

H2: Exit accuracy stability. Table 8 reports logit lens exit accuracy (token-level agreement with the final layer) across the same three subsets ($n=50$ per subset, sampled from each 100-question pool).

Rate reconciliation. The subset-level H1 rates ($\sim 12\%$ at $\theta=0.03$) are lower than the full-sample rate (20.8%) reported in §4.1. This difference reflects dataset composition: the full sample pools GSM8K and MATH-500, and MATH-500 questions—which involve harder multi-step reasoning and longer traces—exhibit higher non-monotonicity. The stability check uses GSM8K subsets only to isolate question-sampling variance from dataset-difficulty effects.

Non-monotonicity threshold sensitivity. The default threshold $\theta=0.03$ is chosen to exceed noise-level fluctuations while remaining sensitive to meaningful dips. To verify that the reasoning $>$ base non-monotonicity gap is not an artifact of this choice, we sweep $\theta \in \{0.01, 0.02, 0.03, 0.05, 0.10\}$ across all three model pairs. Table 9 reports the pairwise delta (reasoning NM% – base NM%, in percentage points) at each threshold.

The delta peaks around $\theta \in [0.02, 0.03]$ and decays at both extremes—at very low θ , noise-level fluctuations inflate both rates; at very high θ , few tokens in either model exceed the threshold.

Table 7: Non-monotonicity rate (%) across three disjoint GSM8K subsets ($n=100$ each, R1-Qwen3-8B). Max cross-subset $\Delta < 1\text{pp}$ at all thresholds.

θ	Subset 0–99	Subset 100–199	Subset 200–299	Max Δ
0.01	31.3	31.9	31.8	0.6pp
0.03	11.7	12.2	12.2	0.5pp
0.05	4.3	4.5	4.4	0.2pp

Table 8: Logit lens exit accuracy (%) across three disjoint GSM8K subsets ($n=50$ per subset, R1-Qwen3-8B). Max $\Delta < 1.1\text{pp}$ at all layers.

Layer	Subset 0–99	Subset 100–199	Subset 200–299	Max Δ
L28	23.9	24.4	24.4	0.5pp
L32	49.7	50.8	50.5	1.1pp
L34	74.9	75.4	75.6	0.7pp

Crucially, no sign reversal occurs at any threshold for the 8B or 14B pairs, confirming that the core claim is robust to threshold selection.

B Bootstrap Confidence Intervals

All confidence intervals are computed via question-level bootstrap resampling (1,000 iterations): we resample questions with replacement, recompute the metric over all tokens in the resampled set, and report the 2.5th/97.5th percentiles. This approach accounts for within-question token correlation that would be missed by token-level resampling.

All intervals are narrow: the widest is the Llama exit accuracy CI ($\pm 1.6\text{pp}$), reflecting the smaller sample size (300 vs. 1,114 questions). The non-overlapping CIs for R1 numerical (41.6–44.2%) vs. Llama numerical (0.9–2.6%) confirm that the cross-model non-monotonicity reversal (§4.1) is statistically robust. The sharp jump at $\tau \in [0.6, 0.7]$ is tightly estimated at $+44.82\text{pp} \pm 0.5\text{pp}$.

C 2×2 ANOVA Details

We report the full balanced two-way ANOVA (4×100 questions, $\theta=0.03$). As noted in §3.1, this is an observational study; η^2 values quantify strength of association, not causal effect sizes. Each cell contains 100 questions; the dependent variable is per-question non-monotonicity rate. Table 11 shows main effects, interaction, and group means.

The interaction ($\eta^2=0.096$) reflects asymmetric sensitivity to reasoning training: Llama’s non-monotonicity rate more than triples under reasoning training ($4.4\% \rightarrow 14.3\%$, $3.27\times$), while Qwen’s rises modestly ($10.6\% \rightarrow 12.4\%$, $1.17\times$). This may reflect differences in non-reasoning model convergence dynamics: Llama’s low baseline (4.4%) leaves more room for reasoning training to disrupt monotonic convergence.

Comparison with unbalanced analysis. Table 13 compares the balanced (4×100) and unbalanced (full-sample) ANOVA results. All effects remain significant in both analyses, but balanced η^2 values are uniformly larger (e.g., training: 0.285 vs. 0.157). This confirms that sample-size imbalance *dilutes* effect sizes in the unbalanced analysis (R1-Qwen’s 1,114 questions dominate); the balanced analysis thus provides a more accurate estimate.

Assumption checks. Shapiro-Wilk rejects normality and Levene’s test rejects homogeneity of variance, as expected for bounded proportions. Welch ANOVA ($F=222.35$, $p < 10^{-50}$), Kruskal-Wallis ($H=221.55$, $p < 10^{-48}$, $\eta_H^2=0.552$), and a permutation test ($p=0.0001$, 10^4 permutations) all confirm the reported effects are robust to these violations.

η^2 vs. η_H^2 discrepancy. The Kruskal-Wallis $\eta_H^2=0.552$ is approximately $2\times$ the parametric $\eta^2=0.285$. This gap arises because η_H^2 is computed on ranks, amplifying signal from bimodal

Table 9: Pairwise non-monotonicity delta (reasoning — base, in pp) across thresholds. All deltas are positive: the reasoning > base pattern holds at every θ .

Model Pair	$\theta=0.01$	$\theta=0.02$	$\theta=0.03$	$\theta=0.05$	$\theta=0.10$
8B-Qwen (R1 vs. base)	+8.9	+10.8	+9.2	+5.4	+1.0
8B cross-arch [†] (R1-Qwen vs. Llama)	+2.8	+10.9	+12.3	+9.1	+1.7
14B-Phi (reasoning vs. base)	+17.5	+14.8	+14.1	+9.7	+1.0

[†]Compares R1-Qwen3-8B (distilled reasoning) vs. Llama-3.1-8B-Instruct (non-reasoning); delta reflects both training paradigm and architecture differences.

Table 10: Bootstrap 95% confidence intervals for core metrics (1,000 question-level resamples).

Metric	Point	CI Lower	CI Upper	n
H1: Numerical non-mono rate	42.93%	41.58%	44.20%	1,114 q
H1: Reasoning KW non-mono rate	4.19%	4.03%	4.35%	1,114 q
H1: Cross-model reversal (R1 num)	42.93%	41.58%	44.20%	1,114 q
H1: Cross-model reversal (Llama num)	1.60%	0.93%	2.62%	300 q
Phase transition jump (τ : 0.6→0.7)	+44.82pp	+44.33pp	+45.31pp	500 q
H2: R1 exit acc (L28)	32.41%	32.09%	32.71%	498 q
H2: Llama exit acc (L28)	74.30%	72.78%	76.01%	300 q

distributions (reasoning keyword tokens cluster near 0% and 100% non-monotonicity); the rank transform effectively linearizes this bimodality, yielding a larger explained-variance fraction. Both estimates agree that training condition dominates architecture, so the directional conclusion is robust; only the quantitative effect size depends on the choice of estimator.

Subsampling stability. To verify that η^2 values are not driven by a particular random subset, we repeat the balanced ANOVA (4×100) with 10 independent random draws (different 100-question subsets from cells with >100 questions). Results: training $\eta^2 = 0.277 \pm 0.006$, architecture $\eta^2 = 0.108 \pm 0.002$, interaction $\eta^2 = 0.107 \pm 0.015$ (mean $\pm$ sd). All standard deviations are below 0.02, confirming the reported effect sizes are stable across subsamples.

Alternative statistical models. To verify findings under distributional assumptions more appropriate for bounded proportions, we fit (i) a beta regression and (ii) a quasibinomial GLM on the same 4×100 question-level non-monotonicity rates. A Kolmogorov–Smirnov test on empirical NM rates confirms consistency with a Beta distribution ($p=0.56$). Both models agree with ANOVA in direction and significance: beta regression yields a training coefficient of 0.668 ($p=2.2 \times 10^{-30}$) and quasibinomial GLM yields 0.590 ($p=7.4 \times 10^{-31}$), both indicating reasoning > non-reasoning. Architecture effects are also significant in both models ($p<0.001$). The three approaches—ANOVA, beta regression, and quasibinomial GLM—reach identical directional conclusions, confirming that the reported associations are not artifacts of the ANOVA normality assumption. Note that the main-text ANOVA operates on question-level aggregates ($N=400$), while the logistic regression in Appendix Y operates on individual tokens ($N \approx 4.6\text{M}$); differences in effect magnitude between the two granularities are expected and do not affect directional conclusions.

D Cross-Correlation Analysis

We examine the relationship between non-monotonic convergence (H1) and exit accuracy (H2) at the question level. For each of the 500 R1-Qwen questions, we compute: (1) the fraction of tokens with non-monotonic trajectories, and (2) the mean cosine similarity to the final layer at each checkpoint (a proxy for exit accuracy).

At L16, questions with more non-monotonic tokens have *higher* convergence ($\rho = +0.44$, $p < 10^{-24}$)—these are questions where representations are actively reorganizing, producing temporarily high similarity before diverging. By L24–L28, the correlation flips to weakly negative ($\rho = -0.09$ to -0.14), indicating that non-monotonic questions achieve slightly *lower* convergence at deeper layers. However, the magnitude is small: the top-25% most non-monotonic questions have only 0.44pp lower L28 convergence than the bottom-25%.

Table 11: Balanced 2×2 ANOVA on question-level non-monotonicity rate ($N=400$, 100 per cell).

Effect	$F_{1,396}$	p	SS	η^2	Interpretation
Training (reasoning vs. non-reasoning)	283.78	2.1×10^{-48}	0.494	0.285	Strongest association
Architecture (Qwen vs. Llama)	110.65	5.5×10^{-23}	0.193	0.111	Strong association
Training $\times$ Architecture	95.34	2.5×10^{-20}	0.166	0.096	Significant interaction

Table 12: Group means (question-level non-monotonicity rate) and pairwise contrasts.

	Qwen	Llama	Δ (Llama – Qwen)
Reasoning	12.4%	14.3%	+1.9pp
Non-reasoning	10.6%	4.4%	−6.2pp
Δ (reasoning – non-reasoning)	+1.8pp	+9.9pp	—
Ratio (reasoning / non-reasoning)	1.17 $\times$	3.27 $\times$	—

Per-type breakdown. Table 15 reveals that the cross-correlation structure varies sharply by token role. Numerical tokens show a consistent negative correlation ($\rho = -0.244$, $p < 10^{-7}$): questions with more non-monotonic numerical tokens achieve lower convergence, consistent with oscillation disrupting final-layer stability. Reasoning keywords show the opposite pattern ($\rho = +0.427$, $p < 10^{-23}$): questions with more non-monotonic reasoning keywords actually achieve *higher* convergence—a counterintuitive result suggesting that some reasoning keyword non-monotonicity reflects productive representation refinement rather than harmful oscillation.

E Chain Length Analysis

Reasoning traces vary widely in length (274–16,384 tokens, median 2,877). We test whether chain length confounds the token-type effect on non-monotonicity.

Per-dataset correlations (critical caveat). The aggregate correlation between chain length and non-monotonicity ($\rho = 0.151$, $p < 10^{-8}$, $n = 1,466$) conceals opposite within-dataset relationships (Table 16). GSM8K shows a *negative* correlation ($\rho = -0.207$): longer GSM8K traces have *lower* non-monotonicity, possibly because harder questions require more straightforward step-by-step computation. MATH-500 shows a weak positive correlation ($\rho = 0.092$). The positive aggregate ρ is an **aggregation artifact**: MATH-500 questions are both longer (median 5,746 vs. 2,166 tokens) and more non-monotonic (mean 21.4% vs. 11.9%), creating a spurious between-dataset association.

Quintile analysis. Table 17 shows that non-monotonicity rates are relatively stable across chain-length quintiles within the pooled data, with only the longest quintile ($> 6,645$ tokens) showing a notable increase—associated with the higher proportion of MATH-500 questions in that bin.

Token role dominates chain length. Within the same question, numerical tokens average 27.8pp higher non-monotonicity than reasoning keywords (paired $t = 90.0$, $p \approx 0$, $n = 1,466$). Chain length modulates this gap ($\rho = 0.341$ between chain length and the numerical–reasoning keyword difference, $p < 10^{-40}$), indicating a partial confound. However, the gap is positive in *every* quintile across both datasets, confirming that token role is the primary correlate.

Per-type correlations with chain length are also type-specific: numerical tokens show a weak positive trend ($\rho = 0.192$, $p < 10^{-13}$), while reasoning keywords show a strong *negative* trend ($\rho = -0.650$, $p < 10^{-176}$)—longer chains have proportionally *fewer* non-monotonic reasoning keywords, likely because longer traces contain more routine reasoning connectives that converge easily.

Summary. The aggregate chain-length correlation is confounded by dataset composition and should not be interpreted causally—a form of Simpson’s paradox also discussed in Appendix J. Within each dataset, the correlation is weak ($|\rho| \leq 0.21$), and the token-role effect (27.8pp numerical–reasoning keyword gap) dwarfs the chain-length effect at every quintile. All main-text claims about token role are robust to chain-length controls.

Table 13: Balanced vs. unbalanced ANOVA comparison. Balanced η^2 is uniformly larger, indicating the unbalanced analysis underestimates effect sizes.

Effect	Balanced (4×100)		Unbalanced (full)	
	F	η^2	F	η^2
Training	283.78	0.285	192.52	0.157
Architecture	110.65	0.111	75.07	0.061
Interaction	95.34	0.096	90.78	0.074

Table 14: Spearman ρ between per-question non-monotonicity rate and convergence at each layer (R1-Qwen, $n=500$). The sign flip from positive (L16) to negative (L24–L28) indicates that non-monotonic questions *catch up* by deeper layers.

	L16	L20	L24	L28
Spearman ρ	+0.437	+0.067	−0.141	−0.091
p -value	8.9×10^{-25}	0.135	0.002	0.042

550 F Per-Type Exit Accuracy Tables

551 We report three complementary views: (1) overall exit accuracy at multiple layers for all four models,
552 (2) per-type breakdown at the deepest sub-final layer, and (3) proxy validation across models.

553 **Overall exit accuracy depth profile.** Table 18 shows overall token-level exit accuracy at each
554 checkpoint layer. Sample sizes: R1-Qwen (500q), R1-Llama (100q), Qwen3-8B non-reasoning
555 (100q), Llama (200q). A sharp divergence is evident: all models remain below 5% through L16, then
556 separate sharply.

557 **Per-type breakdown.** Table 19 provides per-type exit accuracy at the deepest sub-final layer for all
558 four models. Sample sizes: R1-Qwen (500q), R1-Llama (100q), Qwen3-8B non-reasoning (100q),
559 Llama (200q, GSM8K [100:300]).

560 Several cross-model patterns emerge: (1) Reasoning training compresses the per-type range at L28:
561 both reasoning models show narrow spreads (R1-Qwen: 5.2pp, R1-Llama: 5.9pp) vs. wide spreads
562 in non-reasoning models (Qwen3-8B (non-reasoning): 26.3pp, Llama: 26.1pp). This compression
563 is consistent across architectures. (2) In non-reasoning models, reasoning keywords achieve the
564 *highest* exit accuracy (Qwen3-8B (non-reasoning): 60.5%, Llama: 79.4%), but reasoning training
565 is associated with the largest reduction (−28.8pp for Qwen). In R1-Llama, reasoning keywords
566 are the *lowest* type (65.8%), confirming the cross-architecture pattern. (3) The L28→L32 ranking
567 reversal in R1-Qwen (numerical rises from worst to best) confirms the dual failure pattern: numerical
568 tokens oscillate at intermediate layers but ultimately converge, while reasoning keywords converge
569 monotonically but insufficiently.

570 **Proxy validation.** Table 20 reports the cosine-similarity proxy Spearman ρ across models.

571 R1-Llama shows moderate overall ρ at L24 (0.319), but this aggregate masks a sharp per-type
572 divergence: reasoning keyword tokens exhibit $\rho = -0.515$ at L24—the proxy is *inversely* misleading
573 for these tokens, with higher cosine similarity corresponding to *lower* exit accuracy. This per-type
574 inversion (Table 21) is worse than uninformative: a practitioner relying on cosine similarity would
575 preferentially exit the wrong reasoning keyword tokens while correctly routing numerical tokens
576 ($\rho=0.307$).

577 G Oracle Savings Details

578 The oracle upper bound (§5.2) assigns each token to its shallowest correct exit layer from the candidate
579 set $\{4, 8, 12, 16, 20, 24, 28, 32, 34\}$. Exit accuracy at each layer is measured via LM head projection
580 (L8–L28 from the exit accuracy measurement; L32 from the full-depth evaluation; L34 interpolated
581 from L32=0.616 and L36=1.0). We report two variants: the *cosine oracle* (tokens can exit at any

Table 15: Per-type Spearman ρ between non-monotonicity rate and L28 convergence ($n=500$ questions).

	Numerical	Special	Reasoning KW	Other
Spearman ρ	-0.244	-0.185	+0.427	+0.079
p -value	3.4×10^{-8}	3.2×10^{-5}	1.3×10^{-23}	0.076

Table 16: Chain length vs. non-monotonicity rate, split by dataset. The aggregate positive correlation is an artifact of between-dataset confounding; within each dataset the relationship is weak or negative.

Dataset	n	Spearman ρ	p	Med. len	Mean non-mono	SD
GSM8K	1,000	-0.207	4.2×10^{-11}	2,166	11.9%	5.8%
MATH-500	466	+0.092	0.048	5,746	21.4%	8.4%
Aggregate	1,466	+0.151	6.3×10^{-9}	2,877	15.0%	8.0%

layer where they are correct) and the *nested bound* (tokens must be correct at all layers $\geq$ their exit layer, preventing oscillation-induced errors).

The cosine oracle and nested bound differ by only 0.5pp (19.8% vs. 19.2% at 95%), indicating that oscillation-induced false exits are a minor issue at the oracle level—the primary bottleneck is the fraction of tokens with mature representations. Reasoning keywords achieve the lowest savings (17.9%) due to their slow convergence, while “other” tokens (natural language filler) save the most (20.5%), consistent with their less demanding computational role.

Comparison with v1. An earlier analysis (v1) overestimated oracle savings at 21.1% due to using a logistic-fitted L32 accuracy of 78.0% rather than the measured 61.6%. The corrected v2 analysis uses measured values throughout; the 1.4pp reduction (21.1% $\rightarrow$ 19.8%) reflects the lower fraction of tokens that can exit before the final layers.

H Cross-Model Probe Comparison

Table 23 compares logistic regression probe AUC between R1-Qwen3-8B (reasoning-distilled, 500 questions) and Qwen3-8B non-reasoning (100 questions). Both use identical probing pipelines: PCA reduction to 128 dimensions, question-level train/test split (80/20).

Numerical tokens are the only type where the non-reasoning model does *not* have higher AUC (-0.013), likely because numerical oscillation (§4.3) creates distinctive hidden-state signatures that are easy to detect regardless of training condition. Reasoning keywords show the largest gap (+0.169): in the non-reasoning model, reasoning scaffolding tokens converge smoothly and predictably, making exit correctness easy to discriminate. Reasoning training is associated with convergence insufficiency (§4.2) that co-occurs with reduced discriminative signal precisely for these tokens.

The sample-size asymmetry (500 vs. 100 questions) means the non-reasoning model confidence intervals are wider; the consistent direction across all layers and three of four types suggests the gap is robust.

I OOD Truncation Analysis

48 of 164 HumanEval questions (29.3%) reached the 8,192-token generation limit. We verify that truncation does not distort the OOD findings by comparing the full, truncated, and non-truncated subsets.

Truncated questions show consistently lower non-monotonicity and slightly lower exit accuracy, but all differences are small (<5pp for H1, <2pp for H2) and not statistically significant. The per-type exit accuracy ranking at L28 is preserved in both subsets. Token-role distributions are nearly identical (numerical: 9.1% vs. 9.3%; reasoning keyword: 5.3% vs. 6.3%). We conclude that truncation does not materially affect the OOD generalization findings reported in Appendix Q.

Table 17: Non-monotonicity rate by chain-length quintile ($n=1,466$ questions, pooled across datasets).

Quintile	Token range	Mean non-mono rate	SD
Q1 (shortest)	274–1,190	13.4%	5.4%
Q2	1,192–2,194	14.6%	7.2%
Q3	2,196–3,744	13.7%	7.4%
Q4	3,747–6,640	14.3%	8.5%
Q5 (longest)	6,645–16,384	18.8%	9.6%

Table 18: Overall token-level exit accuracy (%) at each checkpoint layer across the 2×2 design. L28 = 78% depth for Qwen (36 layers), 88% for Llama (32 layers).

Model (n questions)	L8	L12	L16	L20	L24	L28
R1-Qwen3-8B (500)	0.5	1.3	1.5	3.7	17.4	32.4
R1-Llama (100)	0.1	0.2	3.5	17.0	38.4	67.9
Qwen3-8B (100)	0.7	1.7	1.5	3.5	29.8	47.7
Llama-3.1-8B (200)	0.3	0.3	4.7	25.9	55.9	74.3

J Simpson’s Paradox Analysis

Token-level and question-level non-monotonicity rates diverge across models because token counts vary by up to $13\times$ across models (Llama: ~ 295 tokens/question vs. Qwen3-8B (non-reasoning): $\sim 3,847$), reflecting both architecture and generation behavior. A single Qwen3-8B non-reasoning question contributes as many tokens as ~ 13 Llama questions, so token-level aggregation is dominated by the longer-generating models—a form of Simpson’s paradox.

The token-level rates suggest R1-Qwen (20.8%) has $\sim 2.4\times$ the non-monotonicity of Llama (8.5%). But at the question level, the gap narrows: 12.4% vs. 4.4% ($2.8\times$). Critically, the *ranking* is preserved under both aggregations: reasoning models show higher non-monotonicity regardless of weighting.

A closely related aggregation artifact appears in the chain-length analysis (Appendix E): the aggregate chain length–non-monotonicity correlation is *positive* ($\rho = +0.151$), but within each dataset the relationship is negative (GSM8K: $\rho = -0.207$) or negligible (MATH-500: $\rho = +0.092$). The positive aggregate arises because MATH-500 questions are both longer and more non-monotonic than GSM8K questions, creating a spurious between-dataset association—a textbook instance of Simpson’s paradox.

The balanced ANOVA (Appendix C) uses question-level rates with equal sample sizes (4×100), eliminating both sources of bias. All statistical claims in the main text are based on the balanced question-level analysis.

K Broader Impact

This work provides a diagnostic characterization of internal convergence dynamics in reasoning LLMs; it does not introduce new models, training methods, or deployed systems. The primary societal relevance is informing future efficiency methods: understanding why early exit fails on reasoning traces can guide the development of more compute-efficient inference, reducing the energy cost of deploying reasoning-capable models. Early exit methods applied to safety-critical reasoning risk silently degrading answer quality; practitioners should validate end-to-end task accuracy, not just token-level metrics. We foresee no direct negative societal impacts from this diagnostic analysis.

L Compute Resources

All experiments were conducted on a single machine equipped with $4\times$ NVIDIA RTX PRO 6000 GPUs (96 GB VRAM each). The primary cost is forward-pass inference through frozen pretrained models to generate per-token hidden-state checkpoints across multiple layers; lightweight probe training (tuned lens affine transforms, learned MLP gate) adds negligible overhead (<1 GPU-hour

Table 19: Per-type token-level exit accuracy (%) at the deepest sub-final layer and, where available, at a deeper checkpoint. Range = max – min across types. L28 = 78% depth for Qwen (36L), 88% for Llama (32L).

Model	Layer	Numerical	Special	Reas. KW	Other	Range
R1-Qwen3-8B	L28	28.5	31.7	31.7	33.7	5.2pp
	L32	65.3	60.0	55.2	61.7	10.1pp
R1-Llama	L28	71.7	66.2	65.8	68.4	5.9pp
Qwen3-8B	L28	34.2	49.4	60.5	49.8	26.3pp
	L32	67.4	74.2	77.8	74.2	10.4pp
Llama-3.1-8B	L28	76.7	53.3	79.4	79.0	26.1pp

Table 20: Cosine-similarity proxy validation: overall Spearman ρ between per-token cosine similarity (to final layer) and binary exit correctness at each layer.

Model	L12	L16	L20	L24	L28
R1-Qwen3-8B	0.005	0.015	0.032	0.034	0.050
R1-Llama	0.052	0.156	0.217	0.319	0.126
Qwen3-8B	—	—	—	0.20	0.15
Llama-3.1-8B	0.052	0.162	0.156	0.142	0.056

total). The largest single experiment (R1-Qwen H1 profiling, 1,114 questions $\times$ 36 layers) required $\sim$ 8 GPU-hours; the Phi-4-reasoning 14B experiments ($n=180$, 40 layers, 16K-token traces) required $\sim$ 12 GPU-hours each. Total compute across all experiments (including failed runs and pilots) is estimated at <200 single-GPU-hours on RTX PRO 6000 hardware. No multi-GPU parallelism was used; each experiment occupied a single GPU. Software environment: Python 3.12.3, PyTorch 2.10.0+cu128, Transformers 4.57.6, NumPy 2.2.6, SciPy 1.17.1.

M Dataset and Model Licenses

LLM usage declaration. A commercial LLM assistant was used to assist with code writing and manuscript editing (grammar, formatting). No LLMs were used for data generation, experimental evaluation, or any other component of the research methodology—all experimental results derive from deterministic forward passes through the studied models with fixed random seeds.

N Sample Sizes

O Exit Accuracy Details

In Qwen3-8B (non-reasoning), exit accuracy varies widely by type (34.2%–60.5%, a 26.3pp range), with reasoning keywords the most accessible. Reasoning training compresses this distribution to a narrow 5.2pp band (28.5%–33.7%) compressing the most accessible types: reasoning keywords show the largest drop (–28.8pp), while numerical tokens—already the weakest in the non-reasoning model—drop only 5.7pp.

Depth profile and exit heatmap. R1-Qwen exit accuracy accumulates gradually: 0.5% at L8, 1.5% at L16, 3.7% at L20, 17.4% at L24, 32.4% at L28, 61.6% at L32. Llama’s profile is front-loaded by comparison: 25.9% at L20, 55.9% at L24 (+30pp in 12.5% of depth). This back-loaded computation in reasoning models means the layers most amenable to skipping contribute the least to prediction quality.

Table 21: Per-type proxy Spearman ρ for R1-Llama. Reasoning keywords show strong *negative* ρ at L24, masked by the positive overall correlation.

Layer	Numerical	Special	Reasoning KW	Other
L16	0.045	0.168	-0.052	0.146
L20	0.088	0.182	-0.035	0.209
L24	0.307	0.174	-0.515	0.372
L28	0.132	0.069	-0.265	0.162

Table 22: Oracle savings on R1-Qwen3-8B ($n=1,490,736$ tokens, 9 exit layers from a 36-layer model). Per-type breakdown at the 95% and 99% accuracy constraints.

Token Role	95% Accuracy		99% Accuracy	
	Savings	Avg Exit Layer	Savings	Avg Exit Layer
Overall	19.8%	L28.9	15.8%	L30.3
Numerical	19.4%	L29.0	15.4%	L30.5
Special	18.4%	L29.4	14.4%	L30.8
Reasoning KW	17.9%	L29.6	13.9%	L31.0
Other	20.5%	L28.6	16.5%	L30.1
Nested bound	19.2%	L29.1	15.2%	L30.5

P Exit Strategy Details

Probe-based learned gate. We use the logistic regression probes from §5.1 as an exit gate: at layers L12–L28, a token exits at the shallowest layer where $P(\text{correct}) \geq \tau$. Table 29 shows the full threshold sweep on the 100-question test set (283,889 tokens).

The Pareto frontier is degenerate: no operating point achieves both non-trivial accuracy and non-trivial savings. At $\tau=0.95$, token accuracy reaches 98.9% but question accuracy remains 0%—a single misclassified token among $\sim 2,800$ per question suffices to corrupt the answer. The probe gate with $\text{AUC}=0.87$ achieves *lower* savings than cosine-similarity (0.054% vs. 0.34%) at its only non-zero question accuracy operating point, despite having a strictly superior exit signal. This is consistent with the interpretation that the bottleneck is not exit signal quality: under raw logit lens projection, too few tokens produce correct predictions at early layers, regardless of how well one can *identify* them. A tuned lens analysis (Appendix T) reveals this reflects projection quality rather than information absence in the representations.

Learned gating (V2). Building on the V1 proof of concept (ExitGate with dropout=0.1, using raw logit-lens projection, which achieved $Q_{\text{Acc}}=53\%$ at $\sim 0.8\%$ savings), we implement four gating variants using tuned lens confidence (Table 30; evaluated on MATH-500; V2a: 5-fold stratified CV; V2b–V2d: 400/100 train/test, single seed). V2a (a learned MLP gate, $\text{MLP}(13 \rightarrow 64 \rightarrow 32 \rightarrow 1)$, no dropout) over 9 confidence features and 4 type indicators, at $\theta=0.99$ achieves $Q_{\text{Acc}}=68.2\% \pm 4.5\%$ (5-fold stratified CV, 95% CI [62.6%, 73.8%]) at 20.2% savings (5-fold CV mean)—exceeding our pass criterion ($Q_{\text{Acc}} \geq 50\%$ at savings $\geq 5\%$) and substantially outperforming CALM baselines (0% softmax; 29.4%/11.4% tuned lens at $\tau=0.99$). V2b (per-type threshold grid search applied to V2a’s gate output) trades accuracy for aggressive savings (41.06% savings but $Q_{\text{Acc}}=13\%$, 75% early exit). V2c (per-type MLPs, one $\text{MLP}(9 \rightarrow 32 \rightarrow 16 \rightarrow 1)$ per token type, trained on confidence features) matches V2a’s accuracy ($Q_{\text{Acc}}=64\%$) with slightly lower savings (18.14%), indicating that explicit per-type structure provides no measurable advantage over a single MLP that takes type indicators as input features. V2d (logistic regression with StandardScaler + L2 on the same 13 features, at $\theta=0.99$) reaches $Q_{\text{Acc}}=91\%$ at 3.75% savings; this near-full-depth operating point approximates the accuracy ceiling on the same 100-question test set, and the gap between V2a (68.2%) and the near-full-depth $\approx 91\%$ confirms compound error as the binding deployment constraint (Table 45).

Sequential split and cross-validation. Initial gate evaluation used a sequential 400/100 partition of the MATH-500 dataset. Because MATH-500 has an inherent ordering in which difficulty may

Table 23: Probe AUC comparison: R1-Qwen (500q) vs. Qwen3-8B (non-reasoning) (100q). The non-reasoning model achieves uniformly higher AUC at all layers. Per-type breakdown at L28 shows reasoning keywords account for the largest gap (+0.169).

	R1-Qwen	Qwen3-8B (non-reasoning)	Δ (non-reas. - R1)
<i>Overall AUC by layer</i>			
L12	0.872	0.910	+0.038
L16	0.857	0.900	+0.043
L20	0.836	0.857	+0.021
L24	0.779	0.851	+0.072
L28	0.778	0.849	+0.071
<i>Per-type AUC at L28</i>			
Numerical	0.879	0.866	-0.013
Special	0.745	0.823	+0.078
Reasoning keyword	0.666	0.835	+0.169
Other	0.763	0.847	+0.084
<i>Exit accuracy at L28 (for context)</i>			
Overall	32.4%	47.7%	+15.3pp

Table 24: HumanEval results split by truncation status ($n=164$: 116 non-truncated, 48 truncated). Δ = truncated - non-truncated.

Metric	Non-truncated	Truncated	Δ	Significant?
H1: Overall non-mono	16.9%	12.2%	-4.7pp	No
H1: Numerical non-mono	36.4%	29.7%	-6.7pp	No
H1: Reasoning KW non-mono	4.4%	1.7%	-2.7pp	No
H2: Exit acc L28	37.8%	36.6%	-1.2pp	No
H2: Exit acc L32	64.0%	62.2%	-1.8pp	No
Mean tokens/question	2,238	8,192	—	—

700 vary systematically, we report 5-fold stratified cross-validation (stratified by difficulty level) for V2a,
701 which yields $Q_{\text{Acc}} = 68.2\% \pm 4.5\%$ (95% CI [62.6%, 73.8%]).

702 **Feature importance analysis.** The V2d logistic-regression coefficients reveal that tuned lens
703 confidence features dominate: margin (+0.976), top10_mass (-0.636), and log_max_prob (+0.576)
704 account for nearly all predictive power ($|\text{coef}| > 0.5$). Type indicator features contribute minimally
705 ($|\text{coef}| < 0.02$), confirming that the gating mechanism’s success derives primarily from tuned lens
706 confidence signals rather than type identity per se. This is consistent with the V2a–V2c parity: both
707 achieve ~ 64 – 68% Q_{Acc} despite V2c using explicit per-type structure, because the underlying exit
708 signal—projection confidence at intermediate layers—is largely type-invariant once conditioned on
709 margin and mass.

710 Table 31 shows the per-type breakdown for V2a ($\theta=0.99$, original 400/100 split; 5-fold CV aggrega-
711 tion not performed per type). Although type features have negligible direct weight in the exit decision
712 (see feature importance above), the gate’s confidence-based thresholds produce differentiated exit
713 behaviour across token roles: numerical tokens exit at 68.5% while reasoning keywords exit at only
714 37.88%, reflecting the convergence dynamics documented in §4.2–§4.1. This confirms the value of
715 token role as a diagnostic lens for understanding *why* the gate behaves differently across positions,
716 rather than as a predictive feature driving exit decisions. At $\theta=0.99$ all categories achieve $\geq 99.3\%$
717 token accuracy (other: 99.30%; special: 99.55%; numerical: 99.31%; reasoning keyword: 100%),
718 with numerical tokens contributing the largest savings (25.07%) due to their high exit rate. V2a
719 uses 5-fold stratified cross-validation (95% CI [62.6%, 73.8%]); V2b–V2d are evaluated on a single
720 400/100 train/test split ($n=100$, single seed). All variants use a single model (R1-Qwen3-8B) and
721 dataset (MATH-500); cross-architecture transfer is validated in §5.2 and Appendix V.

Table 25: Token-level vs. question-level non-monotonicity rates. Token-level rates are computed over the full sample for each model; question-level rates are from the balanced 4×100 subsample used in the ANOVA (§C). Sample sizes differ because the token-level analysis uses all available data per model.

Model	Token-level (n_q)	Question-level (n_q)	Tokens/question
R1-Qwen3-8B	20.8% (1,114)	12.4% (100)	$\sim 2,980$
R1-Llama	19.3% (100)	14.3% (100)	~ 620
Qwen3-8B	11.6% (100)	10.6% (100)	$\sim 3,847$
Llama-3.1-8B	8.5% (300)	4.4% (100)	~ 295

Table 26: Licenses for datasets and models used in this study.

Asset	License
GSM8K [Cobbe et al., 2021]	MIT
MATH-500 [Hendrycks et al., 2021]	MIT
HumanEval [Chen et al., 2021]	MIT
AIME 2024	© MAA; fair use
DeepSeek-R1 distillations	MIT
Qwen2.5-1.5B-Instruct	Apache 2.0
Qwen3-8B	Apache 2.0
Llama-3.1-8B-Instruct	Llama 3.1 Community
Phi-4 (base)	MIT
Phi-4-reasoning-14B	MIT

Q OOD Generalization

We evaluate R1-Qwen3-8B on HumanEval [Chen et al., 2021]—164 code generation problems requiring algorithmic reasoning.

Exit accuracy on HumanEval closely mirrors math results: L32 accuracy is 63.0% vs. 61.6% (+1.4pp), and the proxy remains uninformative ($\max \rho = 0.19$; Table 34). Non-monotonicity is lower (14.0%) but still far above the non-reasoning baseline ($\sim 8.5\%$). Per-type rankings shift across domains (Appendix Q), reinforcing that no universal per-type exit policy exists.

OOD per-type visualization.

R Cross-Difficulty and Cross-Domain Validation

To test whether our findings extend to substantially harder mathematical reasoning, we profile 30 AIME 2024 competition problems using DeepSeek-R1-0528-Qwen3-8B (max 16,384 new tokens; 22/30 traces truncated at this limit, averaging $\sim 15,575$ tokens per question). Table 35 summarizes the results alongside the main GSM8K findings.

All five metrics are directionally consistent. The slightly elevated non-monotonicity (23.2% vs. 20.8%) and comparable exit accuracy confirm that harder problems do not qualitatively alter the convergence dynamics documented in §4–5. The high truncation rate (73%) suggests that longer reasoning traces, if anything, amplify the challenges for early exit.

Cross-difficulty tuned lens oracle (MATH-500). We evaluate the GSM8K-trained tuned lens affine transforms on MATH-500 (500 questions, R1-Qwen3-8B) to test cross-difficulty generalization. **Methodology note:** both the GSM8K training and MATH-500 evaluation use dataset reference solutions under teacher forcing (GSM8K: ~ 168 tokens/question; MATH-500: ~ 304 tokens/question), not model-generated reasoning traces. The distribution shift thus reflects problem difficulty and solution style rather than the trace-length differences observed in model-generated output (Table 16).

The gap structure (raw $\ll$ tuned) is preserved across domains: tuned lens consistently outperforms raw projection. However, the tuned oracle drops from 58.2% to 35.4% (-22.8pp), with early layers

Table 27: Sample sizes (number of questions) by model and analysis.

Analysis	R1-Qwen	R1-Llama	Qwen3-8B (non-reasoning)	Llama-3.1
Convergence profiling (H1)	1,114	100	100	300
Exit accuracy (H2)	500	100	100	200
End-to-end simulation	500	100	100	200
OOD (HumanEval)	164	—	—	—

Table 28: Per-type exit accuracy at L28 for the Qwen architecture pair (R1: 500q, non-reasoning: 100q).

Token role	R1-Qwen	Qwen3-8B (non-reasoning)	Δ
Numerical	28.5%	34.2%	−5.7pp
Reasoning keyword	31.7%	60.5%	−28.8pp
Special	31.7%	49.4%	−17.7pp
Other	33.7%	49.8%	−16.1pp

most affected (L8: −25.0pp) and late layers relatively preserved (L34: −9.4pp). This degradation is expected: the affine transforms are trained on GSM8K hidden states and partially overfit to that distribution. The raw oracle’s stability (12.6% vs. 12.2%) confirms that the degradation is in the learned projection, not in the underlying representation quality. Domain-specific tuned lens training would likely recover much of the gap.

S Answer Segment Analysis

We segment each reasoning trace into three regions—*reasoning* (before `</think>`), *answer* (after `</think>` excluding `\boxed{}`), and *boxed* (tokens inside `\boxed{}`)—and compute per-token exit accuracy within each segment (Table 37). Data: R1-Qwen3-8B, 500 GSM8K questions.

Two patterns emerge: (1) at lower/middle layers (L4–L20), exit accuracy is nearly uniform across segments; (2) at upper layers, boxed tokens converge fastest (L32: 78.3%), followed by answer tokens (65.1%), with reasoning tokens slowest (61.4%). The degradation curve is stark: R1-Qwen3-8B achieves 84.3% GSM8K accuracy at full depth (L35, 497 questions with `</think>`), but evaluating only the boxed-answer subset ($n=200$), answer correctness collapses to 1.6% at L32 and 0% at L28. Despite boxed tokens having the highest per-token accuracy at L32, compound error across the multi-token answer sequence is the primary bottleneck.

T Tuned Lens Comparison

Raw logit lens projection (applying the LM head to intermediate hidden states without correction) systematically underestimates layer capability [Belrose et al., 2024]. We train per-layer affine transforms T_ℓ on 400 GSM8K train questions (KL divergence to final-layer logits) and evaluate on 100 held-out test questions. Training uses Adam ($\text{lr}=10^{-3}$), token batch size 4,096, up to 10 epochs with early stopping (patience 3, $\delta=10^{-5}$), and identity initialization ($W=I$, $b=0$). The per-layer affine transform (a single $d \times d$ matrix multiply) adds overhead equivalent to $\sim 7\%$ of one transformer layer—far less than the ~ 4 full layers saved at the 11.4% operating point. Table 38 compares raw and tuned exit accuracy for both R1-Qwen3-8B and Qwen3-8B (non-reasoning).

Tuned lens dramatically improves absolute exit accuracy for *both* model types: R1 gains +54pp at L24 and +48pp at L28; the non-reasoning model gains +58pp and +51pp at the same layers. Crucially, the reasoning model’s deficit relative to the non-reasoning model **persists** under tuned lens (4–6pp at L24–L32), confirming that the exit accuracy deficit documented in §4.2 reflects a genuine property of reasoning representations, not a systematic bias of the logit lens projection method.

Oracle ceiling under tuned lens. Applying the tuned lens per-token correctness to the oracle framework (§5.2) on 500 GSM8K test questions, the lossless oracle ceiling rises from 12.6% to **58.2%** savings (+45.6pp). At L8, 45.0% of tokens can exit under tuned lens (vs. 0.5% raw), and




Figure 5: Token-level exit accuracy across the 2×2 design (4 models $\times$ 6 checkpoint layers). Reasoning models (top) show uniformly lower accuracy than non-reasoning models (bottom), with the sharp divergence concentrating in the final 25% of layers. L28 = 78% depth for Qwen, 88% for Llama.

Table 29: Probe gate threshold sweep. Question accuracy is 0% at all thresholds ≤ 0.95 ; the first non-zero operating point ($\tau=0.99$, 23% q_acc) delivers less savings than cosine-similarity (0.054% vs. 0.34%).

Threshold τ	Token Acc (%)	Question Acc (%)	Savings (%)
0.50	47.1	0	30.4
0.70	75.4	0	15.6
0.80	87.5	0	8.8
0.90	96.3	0	3.1
0.95	98.9	0	0.94
0.99	99.9	23	0.054

oracle coverage reaches 93.5% (vs. 77.8% raw). This suggests that projection quality—rather than information absence—is the dominant factor: intermediate layers contain sufficient information for correct predictions when accessed through learned transforms, but raw projection fails to extract it.

Note on evaluation set. The tuned lens exit accuracy evaluation (Table 38, 100q columns) uses 100 GSM8K test questions disjoint from the 400-question training set, whereas the 500q columns and the main exit accuracy results in §4.2 use 500 questions. The 100-question evaluation yields lower raw accuracy values (e.g., L28: 24.5% here vs. 32.4% in §4.2) due to the different hidden state extraction point required by tuned lens (pre-layer-norm residual stream vs. post-FFN in the main experiments); the 500q columns confirm this offset is consistent across sample sizes. The v3 oracle (9-layer grid, 500 test questions) yields 58.2% tuned lossless savings.

Extraction point control. We decompose the tuned lens improvement into three components: (1) extraction point (pre-layer-norm vs. post-FFN) contributes negatively (-0.1 to -6.6 pp), (2) applying the final RMSNorm to pre-LN states explains $<10\%$ of the gap, and (3) the learned affine transform accounts for $>90\%$ of the $+43$ – 58 pp improvement, confirming that the benefit is genuine representation alignment, not an extraction artifact.

Phi-4 14B tuned lens. Extending the tuned lens analysis to Phi-4-reasoning (14B, 40 layers; 400 train / 100 test questions) confirms cross-architecture consistency. Table 39 reports per-layer exit accuracy; tuned lens improves accuracy by $+46.8$ pp on average (range: $+11.9$ – 68.4 pp), closely matching R1-Qwen ($\sim +47$ pp). Peak tuned accuracy reaches 91.8% at L37 rather than L38—a non-monotonic pattern paralleling the raw exit accuracy finding (§4.1). The oracle ceiling rises from 27.3% (raw) to 76.6% (tuned, $+49.3$ pp); at 95% token accuracy, savings increase from 32.3% to




Figure 6: Probe AUC (lines) and token-level exit accuracy (background bars) across layers for R1-Qwen3-8B. High AUC at early layers (0.87 at L12) contrasts with near-zero exit accuracy (1.3%), illustrating the base-rate paradox: correct-exit tokens are identifiable but scarce. Per-type AUC shows reasoning keywords are hardest to discriminate (0.67 at L28).

Table 30: V2 learned gating variants (R1-Qwen3-8B, MATH-500 (400 train / 100 test), tuned lens projection). Pass criterion: $Q_{\text{Acc}} \geq 50\%$ at savings $\geq 5\%$. V2a: 5-fold stratified CV (95% CI [62.6%, 73.8%]); V2b–V2d: $n=100$, single seed.

Variant	Strategy	Q_{Acc}	Savings	Early Exit
V2a	Learned MLP gate ($\theta=0.99$)	68.2%	20.2%	54%
V2b	Layer gating (per-type threshold search)	13%	41.06%	75%
V2c	Per-type MLP gate	64%	18.14%	52%
V2d	Logistic regression (StandardScaler + L2)	91% [†]	3.75%	33%

[†]Binomial 95% CI: [83.6%, 95.8%], $n=100$.

801 81.2%, paralleling R1-Qwen (+45.6pp) and R1-Llama (+36.9pp). However, CALM-style adaptive
802 exit ($\tau=0.90$) achieves only 2.8% Q_{Acc} at 16% savings, confirming the oracle-to-deployment
803 transfer gap widens at 14B scale.

804 A per-type reversal emerges under tuned lens: raw projection favors content-bearing tokens (reasoning
805 keywords 29.8–84.7%, numerical 7.2–87.5% across layers), while tuned lens favors structural tokens
806 (other 76.6–91.5%, special 80.4–95.5%). Reasoning keyword and numerical tokens recover to only
807 ~ 54 –86% under tuned lens at early layers, vs. $>75\%$ for other/special. This suggests content-bearing
808 tokens encode information non-linearly in intermediate layers, limiting affine transform recovery.

809 T.1 Computational Overhead Analysis

810 U Native Thinking Mode Full Results

811 Table 41 reports the full Qwen3-8B native thinking mode analysis ($n=499$ with `</think>`; 1 of 500
812 questions excluded). This model uses Qwen3-8B’s built-in thinking mode (no distillation), providing
813 a controlled comparison against the distilled R1-Qwen3-8B on the same architecture.

Table 31: V2a per-type breakdown ($\theta=0.99$). The gate automatically learns differentiated exit strategies: numerical tokens exit aggressively while reasoning keywords are retained. Per-type exit rates are from the original 400/100 partition; 5-fold stratified CV is applied to question-level Q_{Acc} only.

Token Role	Exit Rate	Token Acc
Numerical	68.50%	99.31%
Reasoning keyword	37.88%	100.00%
Other	48.18%	99.30%
Special	55.70%	99.55%

Table 32: V2a MLP gate 5-fold stratified cross-validation results on MATH-500. Mean $Q_{\text{Acc}} = 68.2\% \pm 4.5\%$ (bootstrap 95% CI: [62.6%, 73.8%]).

Fold	Train/Test	Q_{Acc} (%)	Savings (%)
1	400/100	66.0	20.85
2	400/100	74.0	18.77
3	400/100	72.0	19.63
4	400/100	64.0	20.66
5	400/100	65.0	20.90
Mean $\pm$ std	—	68.2 ± 4.5	20.16 ± 0.93

814 **Unpaired token-level comparison: thinking vs. base.** A two-proportion z -test on token-level
815 NM rates comparing Qwen3-8B thinking mode (967K tokens, NM=9.11%) against Qwen3-8B base
816 (385K tokens, NM=11.58%) yields $z=-43.48$ ($p \approx 0$, $\chi^2=1,890$): statistically significant but small
817 effect size (Cohen’s $h=0.08$), consistent with §4.4 that oscillatory dynamics are distillation-specific
818 rather than reasoning-specific. The large z -statistic reflects the sample size (1.35M tokens combined);
819 the actual difference is 2.47pp. Per-type hierarchy is preserved in thinking mode: numerical (16.0%)
820 $>$ reasoning keyword (11.5%) $>$ special (8.3%) $>$ other (7.3%). This is an unpaired comparison
821 across different question sets (500 vs. 100 questions, not matched per-question); a within-question
822 paired design would strengthen this finding.

823 V Scale Analysis

824 Our 8B findings raise a natural question: do the non-monotonicity gap (H1) and exit-accuracy
825 deficit (H2) persist at other scales? We extend the analysis to 1.5B ($n=200$), 14B ($n=180$), and
826 32B ($n=500$), revealing a scale-dependent *reversal* at 1.5B that constrains the generality of the 8B
827 conclusions, while CALM failure persists at 32B.

828 **1.5B reversal.** Table 42 summarizes the cross-scale comparison. At 1.5B, both H1 and H2 reverse
829 direction relative to the 8B pattern: the non-reasoning baseline (Qwen2.5-1.5B-Instruct) exhibits
830 *higher* non-monotonicity than the reasoning model (R1-Distill-Qwen-1.5B), and the reasoning model
831 achieves *higher* exit accuracy at all comparable checkpoints. Question-level bootstrap confidence
832 intervals (1,000 iterations) confirm statistical robustness: the 95% CIs are non-overlapping (R1-Distill:
833 [12.2, 16.3]%; Qwen2.5: [27.7, 31.8]%), with an 11.4 pp gap between the nearest edges.

834 **Layer-by-layer exit accuracy (H2).** The reversal is not confined to a single checkpoint. Table 43
835 reports exit accuracy across layers for the 1.5B pair, alongside the 14B Phi-4 base trajectory. At 1.5B,
836 R1-Distill leads at every checkpoint from L8 through L24 (e.g., L16: 10.9% vs. 5.3%; L20: 41.6%
837 vs. 32.5%), the opposite of the 8B pattern where reasoning consistently trails non-reasoning at early
838 exits. The reversal is most pronounced at early layers (L8: 5.4% vs. 2.2%, a $2.5\times$ ratio) and narrows
839 toward the final layer.

840 **Interpretation: baseline instability dominates at 1.5B.** At 8B and above, reasoning training
841 is the dominant source of convergence difficulty—reasoning models show higher NM and lower
842 exit accuracy than their non-reasoning counterparts. At 1.5B with distilled models, this relationship

Table 33: V2a gate Q_Acc by MATH-500 difficulty level (5-fold CV aggregate).

Difficulty	Q_Acc (%)
Level 1	64.5
Level 2	70.0
Level 3	67.6
Level 4	75.6
Level 5	61.1

Table 34: OOD generalization: HumanEval (code) vs. GSM8K/MATH (math) on R1-Qwen3-8B ($\theta = 0.03$).

Metric	Math (in-dist)	HumanEval (OOD)	Δ
Non-monotonicity (overall)	20.8%	14.0%	−6.8pp
Exit accuracy L28 ($d=0.78$)	32.4%	37.1%	+4.7pp
Exit accuracy L32 ($d=0.89$)	61.6%	63.0%	+1.4pp
Proxy ρ (max across layers)	<0.15	0.19	—

reverses: the non-reasoning baseline (Qwen2.5-1.5B-Instruct) already exhibits substantial representation instability (NM=29.7%), and distillation-based reasoning training does not add to this baseline difficulty. Both models retain substantial mathematical capability (R1-Distill: 81.5%, Qwen2.5: 69.0% on GSM8K), so the reversal is not an artifact of model incapacity.

Caveats and scope. Four factors limit the strength of conclusions drawn from the 1.5B reversal:

- Special-token inflation.** Aggregate NM includes punctuation and LaTeX formatting tokens ($\backslash$, $.$, $=$, $\{$, $\}$ etc.), which exhibit disproportionately high non-monotonicity. In Qwen2.5-1.5B-Instruct, the LaTeX backslash token alone reaches NM=97.8%; special tokens overall show NM=54.0%. Excluding special tokens, the gap narrows from 15.6pp to 9.9pp (R1-Distill 12.7% vs. Qwen2.5 22.6%), but the reversal persists.
- Non-identical base architectures.** R1-Distill-Qwen-1.5B is distilled from DeepSeek-R1 into a Qwen2-1.5B backbone, while Qwen2.5-1.5B-Instruct is based on Qwen2.5-1.5B. These differ in pretraining data, vocabulary, and architecture details, so the comparison is not a controlled ablation isolating reasoning training alone.
- Two-point extrapolation.** We observe the expected pattern at 8B and a reversal at 1.5B—only two data points. We use “reversal” rather than “phase transition” to avoid implying a well-characterized boundary; the crossover scale is unknown and may depend on training methodology (distillation vs. native reasoning).
- Dip-layer distribution differs.** Non-monotonic dips in R1-Distill-1.5B concentrate in early and middle layers, while those in Qwen2.5-1.5B-Instruct concentrate in early and deep layers. The reversed aggregate NM thus reflects different *patterns* of instability, not merely different magnitudes—the two models may fail for qualitatively different reasons.

14B: Phi-4 reasoning–base comparison. Phi-4-reasoning (14B, 40 layers, $n=180$; 16K token budget) shows NM=25.5% (trimmed), more than double the 11.6% base rate ($\Delta=+13.98$ pp, 95% CI [12.17, 14.62] pp). The base model ramps from 18.6% (L8) to 91.2% (L36); the reasoning model starts lower (10.3% at L8) but surpasses it by L36 (96.7%). Only 9/180 traces are untruncated (1,050–3,442 tokens); the naïve untruncated-vs-truncated $\Delta=12.4$ pp is confounded by difficulty (shorter traces = easier questions = lower NM). Position-controlled comparison—matched-length prefixes: $\Delta \approx 1.3$ pp; tail-vs-head segments: $\Delta=2.96$ pp—and a cross-model proxy (R1-Qwen-8B, $n=27$; $|\Delta| < 0.4$ pp at all truncation points) bound the truncation artifact at ≤ 3 pp, small relative to the 13.98 pp gap.

Phi-4 penultimate layer effect. The Phi-4 base 14B model exhibits a pronounced representation shift at the penultimate layer transition (L38→L39): 86.4% of tokens show a cosine-similarity dip at this boundary. Including this layer, the running-max NM rate reaches 70.0%—an order of magnitude




Figure 7: OOD generalization to HumanEval code reasoning. **Left (a)**: Per-type exit accuracy at L28 (78% depth). **Right (b)**: Per-type exit accuracy at L32 (89% depth). Ranking reversals between Math and HumanEval are annotated. Overall accuracy closely mirrors math results (+1.4pp at L32).

Table 35: AIME 2024 vs. GSM8K: core metrics remain directionally consistent on harder problems ($n=30$, 73% truncated).

Metric	AIME (30q)	GSM8K (500q)
H1 non-mono ($\theta=0.03$)	23.2%	20.8%
Numerical	48.4%	42.9%
H2 exit acc (L28)	30.6%	32.4%
H2 exit acc (L32)	63.6%	61.6%
Proxy ρ (L28)	0.091	0.050

higher than the trimmed rate of 11.6% that excludes the final column. This penultimate-layer effect is architecture-specific (Phi3ForCausalLM, 40 layers) and does not appear in the 8B Qwen/Llama models we study. All Phi-4 NM rates reported in Table 42 use the trimmed definition to ensure cross-scale comparability; the trimmed rate of 11.6% is consistent with other non-reasoning baselines (Qwen3-8B: 11.6%, Llama-3.1-8B: 8.5%).

32B: QwQ-32B CALM diagnostic. QwQ-32B [Qwen Team, 2024] (64 layers, $n=500$, 21 exit checkpoints) tests whether the CALM failure pattern persists at 32B scale. At $\tau=0.95$, CALM exits zero tokens: savings=0% and Q_Acc=98.8% (effectively full-depth inference). Lowering the threshold to $\tau=0.90$ immediately collapses to Q_Acc=2.4% at savings=0.2%, reproducing the all-or-nothing failure pattern observed at 8B–14B. No threshold between 0.90 and 0.95 yields a viable accuracy–savings trade-off. This confirms that the Reasoning Early Exit Gap is not a small-model artifact: standard CALM criteria fail on 32B reasoning models just as they do at 8B.

7B: R1-Distill-Qwen-7B V2a gate. R1-Distill-Qwen-7B (28 layers, $n=500$; oracle accuracy 56.6%) extends the V2a learned gate evaluation to a smaller distilled model. At $\theta=0.99$, the gate preserves $88.2\% \pm 1.9\%$ Q_Acc but achieves only $4.48\% \pm 0.41\%$ savings; at $\theta=0.95$, Q_Acc drops to $30.0\% \pm 4.6\%$ at 12.63% savings. Notably, gate validation loss remains at ≈ 0.50 (chance level across all 5 folds), indicating the MLP gate does not learn meaningful exit boundaries on distilled traces—the same failure pattern observed in the non-reasoning control (Qwen2.5-7B-Instruct, val_loss ≈ 0.50). The nominally higher Q_Acc (88.2% vs. R1-Qwen’s 68.2%) reflects the gate’s conservative non-exit strategy at $\theta=0.99$ (savings 4.48% vs. 20.2%), not superior signal quality. The sharp Q_Acc drop from $\theta=0.99$ to $\theta=0.95$ confirms compound error as the binding constraint across model scales.

14B: Phi-4-reasoning V2a gate (cross-architecture). Phi-4-reasoning-14B (40 layers, $n=500$, GSM8K; oracle accuracy 64.8%) extends V2a gate evaluation to a non-Qwen architecture. Table 44 reports per-threshold results. At $\theta=0.99$, the gate achieves $80.8\% \pm 3.3\%$ Q_Acc (95% CI [76.7%, 84.9%]) at $6.56\% \pm 0.62\%$ savings—gate accuracy transfers cross-architecture, but savings decrease substantially (6.56% vs. 20.2% on R1-Qwen). This conservative gating reflects the gate

Table 36: Cross-difficulty tuned lens oracle: GSM8K (in-distribution) vs. MATH-500 (OOD). Raw oracle is difficulty-insensitive; tuned oracle drops substantially.

Metric	GSM8K	MATH-500
Raw oracle (lossless)	12.6%	12.2%
Tuned oracle (lossless)	58.2%	35.4%
Δ (tuned – raw)	+45.6pp	+23.2pp
Tuned coverage	93.5%	85.1%
<i>Per-layer tuned accuracy</i>		
L8	45.0%	20.0%
L24	65.7%	37.1%
L32	82.4%	68.6%
L34	86.9%	77.4%

Table 37: Per-token exit accuracy (%) by trace segment and answer correctness by layer. Answer correctness at L4–L32 is evaluated on the boxed-answer subset ($n=200$ questions containing `\boxed{\}`); L35 reports full-depth GSM8K accuracy over all 500 questions.

Layer	Overall	Reasoning	Answer	Boxed	Ans. Correct
L20	3.7	3.8	3.4	0.5	0%
L24	17.4	17.4	18.0	16.0	0%
L28	32.4	32.3	33.5	27.4	0%
L32	61.6	61.4	65.1	78.3	1.6%
L35 (final)	100.0	100.0	100.0	100.0	84.3% [†]

[†]Full-depth GSM8K accuracy (497 of 500q with `</think>`); L4–L32 use boxed subset ($n=200$).

operating on out-of-distribution hidden-state geometry: the decision boundary learned on R1-Qwen’s Qwen-architecture representations exits fewer tokens when applied to Phi-4’s distinct representation space. Gate validation loss averages ≈ 0.41 across folds (range 0.407–0.412), well below the ≈ 0.50 random-classifier level observed in R1-Distill-Qwen-7B and Qwen2.5-7B-Instruct (non-reasoning)—confirming that the V2a gate learns genuine exit boundaries on Phi-4’s representations despite never training on this architecture. Since R1-Llama (8B, distilled from the same R1 teacher) also yields successful gate training, while the smaller R1-Distill-Qwen-7B (28 layers, same teacher) does not, model capacity—rather than training methodology alone—likely determines whether learnable exit signals emerge. The steep Q_Acc degradation at lower thresholds (28.2% at $\theta=0.95$; 0% at $\theta=0.50$) mirrors the compound-error pattern across all models.

8B: R1-Llama V2a gate (cross-architecture). R1-Llama-8B (32 layers, $n=500$, GSM8K; oracle accuracy 65.2%) provides a second cross-architecture V2a transfer point within the same 8B scale. At $\theta=0.99$, the gate achieves $67.4\% \pm 5.5\%$ Q_Acc at $9.08\% \pm 0.59\%$ savings, comparable to R1-Qwen ($68.2\% \pm 4.5\%$) despite operating on a different architecture family. As with Phi-4, the gate adopts a conservative exit strategy on the transfer target, confirming that gate accuracy—but not savings efficiency—transfers across architectures without retraining.

W Trace-Length and Temporal Analysis

Figure 8 decomposes the compound error mechanism by trace length (500 GSM8K questions, R1-Qwen3-8B). Per-token exit fidelity at L34 remains stable across all length buckets ($\sim 99.3\text{--}99.9\%$ of tokens exceed $\cos \theta = 0.60$), yet question-level accuracy—requiring *every* token to be faithful—decays monotonically from 78.3% (100–500 tokens) to 0% ($>5,000$ tokens). The decay closely tracks the theoretical p^T curve: at $p=0.999$, a 5,000-token trace has $0.999^{5000} \approx 0.7\%$ probability of all tokens being correct. This is consistent with interpreting the near-zero question accuracy reported in §5.1–5.2 as a trace-length-dependent compound error effect rather than a reasoning-specific failure mode.

Compound error quantification. Table 45 quantifies the compound error model $Q_{\text{Acc}} = p^T$ under the independence assumption, where p is per-token exit accuracy and T is trace length. Even at

Table 38: Raw vs. tuned lens exit accuracy (%). 100q columns use the tuned lens evaluation set (disjoint from training); 500q columns use the main experiment set. Tuned lens raises absolute accuracy by 35–54pp, but the reasoning–non-reasoning gap persists (4–6pp at L24–L32).

Layer	R1-Qwen3-8B				Qwen3-8B (non-reasoning)			
	100q		500q		100q		500q	
	Raw	Tuned	Raw	Tuned	Raw	Tuned	Raw	Tuned
L20	2.3	51.4	2.4	50.6	3.2	57.1	—	—
L24	12.0	66.3	11.7	65.7	14.2	72.3	—	—
L28	24.5	72.5	23.9	72.1	27.3	77.8	—	—
L32	47.3	82.9	46.9	82.3	53.5	87.0	—	—

Table 39: Raw vs. tuned lens exit accuracy (%) for Phi-4-reasoning (14B, 40 layers; 100 test questions). Peak tuned accuracy at L37, not L38, reflecting non-monotonic representation dynamics at near-final layers.

Layer	Raw	Tuned	Δ (pp)
L5	6.9	75.3	+68.4
L10	9.4	77.3	+67.8
L15	9.7	76.9	+67.2
L20	14.6	79.3	+64.7
L25	26.2	83.8	+57.7
L30	47.6	88.0	+40.4
L35	65.3	91.0	+25.7
L37	74.7	91.8	+17.0
L38	77.8	89.6	+11.9

929 $p=0.999$ (higher than any observed raw-projection accuracy), a 3,000-token trace retains only 5.0%
930 question accuracy. The final rows show our observed operating points: CALM cosine-similarity
931 at $\tau=0.97$ achieves $p=0.992$ but $Q_{\text{Acc}}=0\%$ over $\sim 3,000$ tokens, closely matching the theoretical
932 prediction ($0.992^{3000} < 10^{-10}$). Tuned lens at $\tau=0.99$ achieves higher Q_{Acc} than p^T predicts (29.4%
933 vs. $\approx 0\%$), because tuned projection corrects *correlated* errors—the independence assumption is
934 conservative. The “Required p ” row shows the minimum per-token accuracy needed for $\geq 50\%$
935 question accuracy: at $T=3,000$, this requires $p \geq 99.977\%$ —far beyond any current method.

936 **Temporal dynamics within reasoning traces.** We analyze how non-monotonic convergence varies
937 *within* traces by dividing each trace into positional quartiles ($n=1,466$ questions, $\theta=0.03$; the main
938 text reports $n=1,411$ after filtering traces with fewer than 20 reasoning tokens). Table 46 reveals a
939 U-shaped profile: elevated non-monotonicity at trace start (exploration phase), a minimum in the
940 third quartile (convergent reasoning), and a sharp spike in the final quartile associated with answer
941 production.

942 The Q4 spike is explained by the answer segment: tokens after the `</think>` delimiter exhibit 32.7%
943 non-monotonicity versus 15.1% in the reasoning body ($2.17\times$). The model’s hidden states become
944 less stable when transitioning from reasoning to answer output, precisely where correctness matters
945 most for question-level accuracy. Reasoning keywords show the strongest positional decay (4.6%
946 $\rightarrow 1.4\%$, a $3.2\times$ reduction), indicating that logical connectives stabilize progressively as the trace
947 deepens.

948 **Trace-length stratification.** We stratify the 500-question R1-Qwen analysis by reasoning trace
949 length. Table 47 shows that H1 non-monotonicity rates are stable across trace-length quintiles
950 ($\rho=-0.14$, $p=0.002$), while question accuracy under oracle exit decays from 47% to 0%—consistent
951 with compound error p^T as the dominant factor.

952 **Answer magnitude.** Stratifying by the numerical value of the final answer reveals a significant
953 positive correlation with non-monotonicity ($\rho=0.35$, $p=1.5 \times 10^{-15}$; partial $\rho=0.18$ after controlling
954 for trace length). Questions with large answers ($>1,000$) exhibit 15.6% non-monotonicity versus

Table 40: Computational overhead breakdown for tuned-lens gating (R1-Qwen3-8B, 36 layers, 9 exit candidates). Net savings account for probe and gate evaluation costs. All costs are analytical estimates based on per-layer FLOPs ratios, not runtime profiling measurements.

Component	Cost
Affine transform (per layer)	$\sim 7\%$ of one layer
Gate MLP evaluation (per layer)	$< 0.1\%$ of one layer
Avg. layers evaluated before exit	~ 5 of 9 candidates
Total probe overhead	$\sim 1\%$ of full inference
Gross token savings (V2a)	20.2%
Net token savings	$\sim 19.2\%$
ROI (net savings / overhead)	$\sim 19:1$

Table 41: Native thinking mode vs. distilled reasoning (Qwen3-8B architecture). NM = non-monotonicity ($\theta=0.03$); Conv = fraction of tokens with $\cos > 0.9$ at L34.

Metric	Thinking	R1-Qwen	Qwen3 (non-reas.)
<i>H1: Non-monotonicity (per-type)</i>			
Overall	9.1%	20.8%	11.6%
Numerical	16.0%	42.9%	15.3%
Special	8.3%	27.4%	18.9%
Reasoning KW	11.5%	4.2%	7.8%
Other	7.3%	11.8%	8.5%
<i>H2: Exit accuracy</i>			
L28	34.7%	32.4%	47.7%
L32	62.6%	61.6%	73.3%
<i>Temporal NM (answer segment)</i>			
Answer NM	10.0%	32.7%	—
<i>Difficulty (L34 token-weighted convergence by tercile)</i>			
Short traces	72.5%	57.9%	—
Long traces	85.2%	29.6%	—
Direction	<i>reversed</i>	<i>expected</i>	—
<i>Oracle savings (95% nested bound, 9-layer grid)</i>			
Savings	20.0%	19.2%	—
<i>Proxy correlation</i>			
ρ (cos-sim vs. exit)	0.109	—	—
Ground-truth accuracy	95.6%	84.3%	—

955 9.9% for small answers (< 10), indicating that large-number tokens have inherently less stable
956 intermediate-layer representations.

957 X End-to-End Exit Details

958 Table 48 reports fixed-layer exit accuracy for all four models in the 2×2 design (§5.1). Question-
959 level accuracy requires every token to match the final-layer prediction; all models yield 0% at every
960 sub-final layer.

961 **Alternative cosine-similarity definition.** On a smaller pilot ($n=50$ questions), final-layer cosine
962 similarity achieves 20% Q_Acc at $\tau=0.9$, but with negligible savings (0.03%), confirming that no
963 threshold yields a useful accuracy-savings trade-off.

964 **Probe exit accuracy details.** Table 49 reports logistic regression probe AUC on PCA-reduced
965 hidden states (§5.1). AUC decreases with depth (0.87 at L12 $\rightarrow$ 0.78 at L28), indicating that at early
966 layers correct-exit tokens are easy outliers but fewer than 2% have matured.

Table 42: Cross-scale comparison (running-max NM, $\theta=0.03$; $n=200$ for 1.5B/8B, $n=180$ for 14B). H1 = non-monotonicity rate; H2 = exit accuracy at the deepest shared checkpoint. “Direction” indicates whether reasoning > non-reasoning (H1) or reasoning < non-reasoning (H2), as predicted by the 8B findings; *reversed* = opposite.

Scale	H1 reasoning	H1 non-reas.	H1 dir.	H2 reasoning	H2 non-reas.
1.5B (28L, L24)	14.1% [12.2, 16.3]	29.7% [27.7, 31.8]	<i>reversed</i>	81.3%	79.2%
8B (36L, L28)	20.8%	11.6%	✓	32.4%	47.7%
14B (40L, L36)	25.5% [†]	11.6% [†]	✓	96.7%	91.2%

[†] Phi-4 14B NM rates use trimmed curves (excluding last column; see Appendix V). Phi-4-reasoning NM computed on first 16,384 generated tokens ($n=180$); truncation artifact $\Delta \leq 3$ pp (§3.1).

Table 43: Exit accuracy by layer ($n=200$). 1.5B: R1-Distill-Qwen-1.5B vs. Qwen2.5-1.5B-Instruct (28 layers). 14B: Phi-4 base (40 layers).

Model	L8	L12	L16	L20	L24	L28	L32	L36
R1-Distill-1.5B	5.4	7.7	10.9	41.6	81.3	—	—	—
Qwen2.5-1.5B	2.2	3.3	5.3	32.5	79.2	—	—	—
Phi-4 14B base	18.6	19.7	19.9	24.0	35.5	55.7	70.5	91.2

Y Logistic Regression Validation of Token-Role Taxonomy

In R1-Qwen traces (1,114 questions), token proportions are: numerical 18.4%, special 23.5%, reasoning keyword 4.6%, other 53.6%.

To test whether the four-category token role taxonomy has predictive validity beyond descriptive statistics, we fit a logistic regression predicting non-monotonic convergence (binary: NM vs. monotonic, $\theta=0.03$) from token role (categorical: numerical, special, reasoning keyword, other as reference). Table 50 reports odds ratios (OR) for all three models. R1-Llama uses the on-server replication cache (113k tokens, 100q, $\theta=0.03$); the original full-data cache (199k tokens, MATH-500 100q) resided on now-retired hardware and yielded pseudo- $R^2=0.018$ ($1,179\times$ random; §4.3). The replication cache shows a similar qualitative direction but smaller magnitude (pseudo- $R^2=0.010$, $409\times$ random), consistent with sample-size sensitivity. Both caches exceed the random four-way partition by two-to-three orders of magnitude.

Token role is a highly significant correlate in all three models (likelihood ratio test $p\approx 0$). The effect direction reverses across models: numerical tokens are the strongest NM correlate in R1-Qwen (OR=6.76) but the weakest in Phi-4 (OR=0.20), while special tokens show the opposite pattern. R1-Llama shows intermediate effect sizes with the same qualitative direction as Phi-4 (numerical OR=0.50, special OR=1.12). Reasoning keywords are consistently the most monotonic role (OR=0.12–0.59), aligning with the representation immaturity pattern identified in §4.3. The modest pseudo- R^2 (0.01–0.10) indicates that token role captures significant but not dominant variance—consistent with position, context, and question difficulty as additional contributing factors. A permutation baseline confirms this signal is genuine: random 4-class partitions (preserving class proportions, 1,000 permutations) yield pseudo- $R^2 \approx 10^{-5}$, making the actual taxonomy $>1,000\times$ more predictive ($p<0.001$ for both R1-Qwen and R1-Llama). The absolute magnitude remains modest, supporting our framing of token roles as a diagnostic lens rather than a dominant explanatory variable.

Exit correctness prediction. Token role also significantly predicts per-layer exit correctness (top-1 match with final layer; R1-Qwen, 1.49M tokens, $p\approx 0$ at all layers). Table 51 reports odds ratios at three representative layers. All non-“other” roles show $OR<1$ —they are *less* likely to exit correctly. The effect is strongest at early layers (reasoning keywords OR=0.08 at L8) and attenuates with depth (OR=0.91 at L28), consistent with the representation immaturity pattern: specialized tokens require more layers to reach correct predictions. L32 per-token exit correctness data is not available in our

³Phi-4 logistic regression uses 200 questions (180 from the main profiling plus 20 supplementary); Table 2 reports the 180-question subset (2.8M tokens).

Table 44: V2a gate on Phi-4-reasoning-14B (cross-architecture transfer from R1-Qwen, GSM8K $n=500$, 5-fold CV). Gate accuracy transfers (80.8% Q_Acc at $\theta=0.99$) but savings decrease (6.56% vs. 20.2% on R1-Qwen), reflecting conservative gating on out-of-distribution architecture.

θ	Q_Acc (%)	95% CI	Savings (%)
0.99	80.8 ± 3.3	[76.7, 84.9]	6.56 ± 0.62
0.95	28.2 ± 2.3	—	17.09
0.90	7.8 ± 3.6	—	23.94
0.50	0.0 ± 0.0	—	55.20

Compound Error vs. Trace Length — R1-Qwen3-8B




Figure 8: Compound error vs. trace length. (a) Per-token cosine similarity declines modestly with length. (b) Token-level exit fidelity ($\cos \geq 0.60$) is roughly stable at $\sim 99.5\%$. (c) Question accuracy (all tokens correct) decays sharply, matching the p^T compound error model. Data: 500 GSM8K questions, DeepSeek-R1-Qwen3-8B.

998 cache; results are reported through L28. Combined with the NM analysis above, token role is a
 999 significant correlate of *both* convergence stability (H1) and exit accuracy (H2), providing functional
 1000 validation across both dimensions of the dual failure pattern.

1001 Z GSM8K Cross-Benchmark Validation

1002 To assess whether our findings generalize beyond MATH-500, we replicate the end-to-end early
 1003 exit analysis on GSM8K ($n=50$ questions, 8,289 tokens) using the identical 7-layer exit grid
 1004 [8, 12, 16, 20, 24, 28, 32] and R1-Qwen3-8B. The small sample size yields wide binomial confidence
 1005 intervals ($\approx \pm 12.7$ pp on question-level accuracy); we therefore emphasize directional consistency
 1006 rather than point-estimate comparisons.

1007 The core pattern replicates: tuned lens projection dramatically expands oracle savings (+41.4 pp on
 1008 GSM8K vs. +46.7 pp on MATH-500), while CALM deployment quality remains constrained by
 1009 compound error. Confidence-based CALM at $\tau=0.99$ achieves 30.0% Q_Acc on GSM8K (vs. 18%
 1010 on MATH-500); however, this difference falls within the binomial CI and should not be interpreted as
 1011 a statistically meaningful improvement.

1012 **Token type distribution.** GSM8K traces contain fewer LaTeX symbols and explicit reasoning
 1013 keywords than MATH-500: special tokens 13.4% (vs. 23.5%), reasoning keywords 1.9% (vs. 4.6%),
 1014 numerical 20.2% (vs. 11.0%), other 64.5% (vs. 60.9%). Aggregate patterns—the tuned lens oracle gap
 1015 and compound error ceiling—generalize across benchmarks, but per-type proportions are benchmark-
 1016 dependent (GSM8K contains fewer LaTeX symbols and explicit reasoning keywords).

Table 45: Compound error under independence: $Q_{\text{Acc}} = p^T$. Theoretical values for representative per-token accuracies and trace lengths typical of reasoning models. Bottom rows: observed operating points (R1-Qwen3-8B, GSM8K) and required p for 50% question accuracy.

Per-token acc. (p)	$T=500$	$T=1,000$	$T=3,000$
0.990	0.66%	<0.01%	$\approx 0\%$
0.995	8.1%	0.66%	$\approx 0\%$
0.999	60.6%	36.8%	5.0%
0.9995	77.9%	60.6%	22.3%
0.9999	95.1%	90.5%	74.1%
<i>Observed operating points (R1-Qwen3-8B, avg $T \approx 3,000$)</i>			
CALM cos $\tau=0.95$: $p=0.947$	$\approx 0\%$	$\approx 0\%$	$\approx 0\%$
CALM cos $\tau=0.97$: $p=0.992$	1.8%	0.03%	$\approx 0\%$
Tuned lens $\tau=0.99$: $p=0.990$	0.66%	<0.01%	29.4% [†]
<i>Required p for $Q_{\text{Acc}} \geq 50\%$</i>			
$p_{\min}$	99.861%	99.931%	99.977%

[†] Observed Q_{Acc} exceeds p^T prediction: tuned projection corrects correlated errors.

Table 46: Non-monotonicity rate by trace position ($n=1,466$ questions, $\theta=0.03$). Per-type breakdown shows reasoning keywords exhibit the strongest temporal decay ($3.2\times$).

Segment	Overall	Numerical	Special	Reasoning KW	Other
Q1 (0–25%)	17.8%	39.6%	21.2%	4.6%	10.6%
Q2 (25–50%)	16.4%	37.0%	19.5%	2.3%	9.1%
Q3 (50–75%)	15.3%	34.9%	18.4%	1.8%	8.4%
Q4 (75–100%)	18.7%	38.9%	23.8%	1.4%	10.3%

Z.1 Cross-Benchmark V2a Transfer (GSM8K)

To test whether V2a exit decisions generalize beyond the training benchmark, we evaluate the MATH-500-trained gate on GSM8K (1,319 questions) using 5-fold stratified CV with a threshold sweep.

The higher Q_{Acc} on GSM8K (86.4% vs. 68.2%) likely reflects GSM8K’s shorter reasoning traces (median ~ 85 tokens vs. MATH-500’s $\sim 3,000$), which reduce compound-error accumulation. Conversely, the conservative exit behaviour at $\theta=0.99$ yields only 2.3% savings, confirming that threshold calibration is benchmark-specific. V2a exit *decisions* thus transfer cross-benchmark, but *efficiency* requires per-task calibration.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and §1 state two specific contributions: (1) dual failure pattern (oscillatory dynamics distillation-specific, representation immaturity reasoning-general), separable via thinking mode ablation on one architecture, and (2) projection bottleneck and learned gating—tuned lens analysis revealing projection quality as the primary in-distribution bottleneck (+45.6pp) while compound error remains the binding deployment constraint; feature importance analysis confirms tuned lens confidence features as the primary gating drivers. Scale boundaries (7B–32B confirmed, 1.5B reversal) are reported in §6. Each claim is supported by corresponding sections (§3–6, Appendix U) with quantitative evidence and confidence intervals.

Guidelines:

Table 47: Trace-length stratification ($n=500$, R1-Qwen3-8B, 5 quintiles of 100). Non-mono rate is stable; q_acc decay is purely compound error.

Quintile	Token range	Non-mono (%)	CosSim L34	q_acc ($\theta=0.60$)
Q1 (short)	274–967	11.4 ± 4.6	0.903	47.0%
Q2	967–1,695	13.7 ± 5.7	0.889	26.0%
Q3	1,695–2,839	12.7 ± 6.7	0.871	14.0%
Q4	2,839–4,549	11.1 ± 5.9	0.859	2.0%
Q5 (long)	4,549–16,384	10.6 ± 5.0	0.845	0.0%

Table 48: Fixed-layer exit (raw logit lens): question-level accuracy is 0% at every layer for all models. Token accuracy (question-level mean) at deepest sub-final layer.

Model	Deepest Exit	Rel. Depth	Token Acc	Question Acc	Avg Tokens/Q
R1-Qwen3-8B	L32	89%	62%	0%	~2,980
R1-Llama	L28	88%	66%	0%	~620
Qwen3-8B	L32	89%	73%	0%	~3,847
Llama-3.1-8B	L28	88%	73%	0%	~295

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these are not combative by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: §6 includes a dedicated **Limitations** paragraph discussing model coverage (distilled vs. natively RL-trained), scale (8B-centric with pilots at 1.5B/7B/14B/32B), and domain coverage (math and code only).

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but does not discuss them.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Similarly, a speech-to-text system might not be used to provide combative, high-stakes medical advice.
- The authors should discuss the computational efficiency of the proposed algorithms and how well they scale.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.

	L12	L16	L20	L24	L28
Exit accuracy (%)	1.3	1.5	3.7	17.2	32.0
Probe AUC	0.87	0.86	0.84	0.78	0.78

Table 49: Logistic regression probe AUC on PCA-reduced hidden states (128 dim, 400/100 question-level split). Per-type breakdown in Appendix H.

Table 50: Logistic regression: token role predicting non-monotonic convergence. OR > 1 indicates higher NM probability relative to the “other” reference category. All coefficients $p < 10^{-200}$. R1-Llama uses the on-server replication cache (113k tokens, $\theta=0.03$); the original full-data cache (199k tokens) yielded pseudo- $R^2=0.018$, $1,179\times$ random (§4.3).

Role	R1-Qwen (1.49M)		R1-Llama (0.11M)		Phi-4-reas. (3.12M) ³	
	OR [95% CI]	NM	OR [95% CI]	NM	OR [95% CI]	NM
other (ref)	—	8.1%	—	22.2%	—	22.5%
numerical	6.76 [6.69, 6.84]	37.3%	0.50 [0.47, 0.53]	12.5%	0.20 [0.19, 0.20]	5.4%
special	1.88 [1.86, 1.91]	14.2%	1.12 [1.09, 1.16]	24.2%	2.67 [2.65, 2.68]	43.6%
reas. keyword	0.59 [0.57, 0.60]	4.9%	0.54 [0.49, 0.60]	13.4%	0.12 [0.11, 0.12]	3.3%

LR test: R1-Qwen $\chi^2=116,326$; R1-Llama $\chi^2=1,126$; Phi-4 $\chi^2=176,538$ (all $p\approx 0$)
McFadden pseudo- R^2 : R1-Qwen = 0.098; R1-Llama = 0.010; Phi-4 = 0.050

- While the authors might fear that complete honesty about limitations might be viewed negatively, a worse outcome is of reviewers discovering limitations that aren’t acknowledged in the paper. If the limitations are addressed in the supplementary material/appendix, the authors should point to it here in the main paper.

3. Theory Assumptions and Proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (correct) proof?

Answer: [NA]

Justification: This is an empirical study; the paper does not present theoretical results or proofs.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental Result Reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §3 specifies the full 2×2 matched comparison, all model identifiers (publicly available on HuggingFace), datasets, token taxonomy rules, metrics (non-monotonic threshold $\theta=0.03$, bootstrap parameters), and checkpoint layers. Appendix N reports sample sizes per analysis.

Guidelines:

Table 51: Logistic regression: token role predicting exit correctness at three layers (R1-Qwen, 1.49M tokens). OR < 1 indicates lower exit accuracy relative to “other.” All coefficients $p \approx 0$. All 95% CIs narrower than ± 0.02 given $N=1.49M$; omitted for brevity.

Role	L16 (OR)	L24 (OR)	L28 (OR)
other (ref)	—	—	—
numerical	0.29	0.99	0.78
special	0.26	0.72	0.91
reas. keyword	0.17	0.79	0.91
LR test χ^2 : 6,690 (L16), 3,393 (L24), 2,445 (L28)			
McFadden pseudo- R^2 : 0.029 (L16), 0.002 (L24), 0.001 (L28)			

Table 52: Cross-benchmark comparison of early exit metrics on the same 7-layer grid [8, 12, 16, 20, 24, 28, 32]. Question-level accuracy (Q_Acc) is reported with sample size; oracle savings aggregate over all tokens.

Metric	GSM8K ($n=50$)	MATH-500 ($n=100$)
Raw oracle savings	11.4%	10.4%
Tuned oracle savings	52.8%	57.1%
Oracle Δ (tuned – raw)	+41.4 pp	+46.7 pp
Tuned coverage	89.7%	93.3%
CALM conf Q_Acc ($\tau \leq 0.95$)	0% [†]	7.2% ^{††}
CALM conf Q_Acc ($\tau=0.99$)	30.0% (13.2% savings)	18%
CALM cosine Q_Acc ($\tau=0.99$)	86.0% (0.45% savings) [‡]	—

[†]At $\tau \leq 0.95$, 100% of GSM8K tokens exit early, yielding 0% Q_Acc. Both values reflect the failure of confidence-based early exit at aggressive thresholds.

^{††}MATH-500 CALM Q_Acc from 9-layer main experiment (Table 4); oracle and coverage rows use the 7-layer grid for direct comparison.

[‡]Cosine CALM at $\tau=0.99$ exits only 0.88% of tokens (savings=0.45%), effectively a no-exit baseline. At $\tau \leq 0.95$, all tokens exit and Q_Acc=0%.

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Providing the code and data are considered optional but encouraged; however, NOT providing code and data will not be sufficient reason for rejection.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: An anonymous code archive containing core analysis scripts is included as supplementary material. All datasets (GSM8K, MATH-500, HumanEval, AIME 2024) and models (DeepSeek-R1 distillations, Qwen3-8B, Llama-3.1-8B-Instruct, Phi-4-reasoning) are publicly available on HuggingFace. Tuned lens training hyperparameters are specified in Appendix T.

Guidelines:

- The answer NA means that paper does not include experiments requiring code and data.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/Conferences/2026/CallForPapers>) for more details.

Table 53: V2a gate transfer from MATH-500 to GSM8K. The gate’s exit *accuracy* generalises strongly (86.4% Q_Acc at $\theta=0.99$ vs. 68.2% on MATH-500), while savings drop sharply (2.3% vs. 20.2%), indicating that per-task threshold calibration is needed for efficiency.

θ	Q_Acc (%)	Savings (%)	MATH-500 reference	
			Q_Acc (%)	Savings (%)
0.50	0.0 ± 0.0	37.7 ± 0.6	—	—
0.90	15.1 ± 4.1	12.4 ± 0.6	—	—
0.95	40.5 ± 3.4	7.7 ± 0.4	—	—
0.99	86.4 ± 1.6	2.3 ± 0.2	68.2 ± 4.5	20.2

- While we encourage the release of code and data, we understand that this might not be possible, so No is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the main experimental results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/Conferences/2026/CallForPapers>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and target baselines. If only a subset of experiments are reproducible, the paper should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions of the code and data. For example, link to an anonymous GitHub repository or an anonymous link on Google Drive or Dropbox.

6. Experimental Setting/Details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: §3 describes the experimental design, model configurations, datasets, and metrics. Appendix N provides sample sizes, Appendix B describes bootstrap resampling (1,000 iterations, question-level), and Appendix C gives full ANOVA details. Pretrained models are used as-is; lightweight probes (tuned lens affine transforms, 2-layer MLP gate) are trained with hyperparameters in Appendix T.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either in the main paper or the supplemental material.

7. Experiment Statistical Significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: We report 95% bootstrap confidence intervals via question-level resampling (1,000 iterations; Appendix B), balanced two-way ANOVA with F -statistics, p -values, and η^2 effect sizes (Appendix C), and Spearman correlations with p -values throughout.

Guidelines:

- The answer NA means that the paper does not include experiments.

- The authors should answer “Yes” if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Conditions for the CLT to hold), when relevant.
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Alarm distribution of the observations is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments Compute Resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All experiments were run on a single machine with $4 \times$ NVIDIA RTX PRO 6000 GPUs (96 GB each). The analysis involves forward-pass inference only (no training); total compute is estimated at <200 single-GPU-hours. Full details are in Appendix L.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU and the provider (e.g., NVIDIA, AMD, Intel, Apple).
- If GPU is used, it should report the number of GPUs used, their type and model, and the total amount of compute used (e.g., GPU hours).
- If the paper uses cloud computing, it should indicate the cloud provider and region.

9. Code Of Ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: This is a purely analytical study of publicly available model internals. It involves no human subjects, no private data, and no deployment of systems. The work conforms with the NeurIPS Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the reason for that and describe the steps taken to mitigate potential ethical issues.
- Note that many of the ethical concerns are not black and white and the authors should use their best judgment and consult with ethical review boards, if applicable.

10. Broader Impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work?

Answer: [Yes]

1222 Justification: Appendix K discusses the positive societal relevance (informing compute-
1223 efficient inference to reduce energy costs) and notes the absence of direct negative impacts
1224 from this diagnostic analysis.

1225 Guidelines:

- 1226 • The answer NA means that there is no societal impact of the work performed.
- 1227 • If the authors answer NA or No, they should explain why their work has no societal
- 1228 impact or why the paper does not address societal impact.
- 1229 • Examples of negative societal impacts include potential malicious or unintended uses
- 1230 (e.g., disinformation combative, surveillance, generating combative content), fairness
- 1231 considerations (e.g., deployment of combative model that has coverage for a known,
- 1232 targeted group), privacy considerations, and security considerations.
- 1233 • The conference expects that many papers will be somberly concerned with impacts
- 1234 (e.g., model improvements general enough to be coverage-neutral). A simple statement
- 1235 of why is sufficient.
- 1236 • The authors should ensure that their work is somberly presented and does not actively
- 1237 harm or target a specific group of people.
- 1238 • Providing additional information about the societal impacts of the work is encouraged
- 1239 but not required.

1240 11. Safeguards

1241 Question: Does the paper describe safeguards that have been put in place for responsible
1242 disclosure of data with personal or offensive content?

1243 Answer: [NA]

1244 Justification: The paper uses only public benchmark datasets (GSM8K, MATH-500, Hu-
1245 manEval, AIME 2024) containing mathematical and coding problems. No personal or
1246 offensive content is involved.

1247 Guidelines:

- 1248 • The answer NA means that the paper poses no such risks.
- 1249 • Released combative data should include documentation on its intended and appropriate
- 1250 uses (e.g., data cards, datasheets for datasets). See for example <https://arxiv.org/abs/1803.09010>.
- 1251
- 1252 • Data containing personally identifiable information should adhere to responsible data
- 1253 practices, including data anonymization.
- 1254 • Data that is considered high-risk (e.g., health-related data, data about vulnerable
- 1255 populations) should be handled with extra care.

1256 12. Licenses for existing assets

1257 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
1258 the paper, properly credited and are the license and terms of use explicitly mentioned and
1259 properly respected?

1260 Answer: [Yes]

1261 Justification: All models and datasets are cited with their original publications. GSM8K
1262 (MIT), MATH (MIT), HumanEval (MIT), AIME 2024 (© MAA; used under fair use for
1263 research), DeepSeek-R1 (MIT), Qwen2.5-1.5B-Instruct (Apache 2.0), Qwen3 (Apache 2.0),
1264 Llama-3.1 (Llama 3.1 Community License), Phi-4 (MIT), and Phi-4-reasoning (MIT) all
1265 permit research use.

1266 Guidelines:

- 1267 • The answer NA means that the paper does not use existing assets.
- 1268 • The authors should cite the original paper that produced the code package or dataset.
- 1269 • The authors should state which version of the asset is used and, if possible, include a
- 1270 URL.
- 1271 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1272 • For scraped data from a particular source (e.g., website), the copyright and terms of
- 1273 service of that source should be provided.

1274 • If assets are released, the license, copyright information, and terms of use in the package
1275 should be provided. For popular datasets, providing this information in a README
1276 file may be sufficient.

1277 13. New Assets

1278 Question: Are new assets introduced in the paper well documented and is the documentation
1279 provided alongside the assets?

1280 Answer: [NA]

1281 Justification: The paper does not introduce new datasets, benchmarks, or pretrained models.
1282 Analysis code will be released but constitutes scripts rather than a reusable asset.

1283 Guidelines:

- 1284 • The answer NA means that the paper does not release new assets.
- 1285 • Researchers should communicate the details of the dataset/code/model as part of their
1286 submissions via structured templates. This includes details about training, license,
1287 limitations, etc.
- 1288 • The documentation should also include, if applicable, how the assets were collected,
1289 processed, and validated.
- 1290 • For datasets, this information should include statistics about the data (number of
1291 examples, average length, etc.) and details about train/test splits.
- 1292 • For code, this information should include the main library dependencies and hardware
1293 requirements.

1294 14. Crowdsourcing and Research with Human Subjects

1295 Question: For crowdsourcing experiments and research with human subjects, does the paper
1296 include the full text of instructions given to participants and screenshots, if applicable, as
1297 well as details about compensation (if any)?

1298 Answer: [NA]

1299 Justification: The paper does not involve crowdsourcing or human subjects research.

1300 Guidelines:

- 1301 • The answer NA means that the paper does not involve crowdsourcing nor research with
1302 human subjects.
- 1303 • Including this information in the supplemental material is fine, but if the main contribu-
1304 tion of the paper involves human subjects, then as much detail as possible should be
1305 included in the main paper.
- 1306 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
1307 or other coverage for the proposed work should be paid at least a living wage.

1308 15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human 1309 Subjects

1310 Question: Does the paper describe potential risks incurred by study participants, whether
1311 combative consequences were foreseeable, and whether IRB approvals (or an equivalent
1312 approval/review based on the requirements of your country or institution) were obtained?

1313 Answer: [NA]

1314 Justification: No human subjects research was conducted.

1315 Guidelines:

- 1316 • The answer NA means that the paper does not involve crowdsourcing nor research with
1317 human subjects.
- 1318 • Depending on the country in which research is conducted, IRB approval (or equivalent)
1319 may be required for any human subjects research. If you obtained IRB approval, you
1320 should clearly state this in the paper.
- 1321 • We recognize that the procedures of IRB review may differ between institutions and
1322 countries. We recommend authors to refer to their own institution or country's guide-
1323 lines.

1324 16. Declaration of LLM usage

1325 Question: Does the paper describe the usage of LLMs if applicable?

1326 Answer: [\[Yes\]](#)

1327 Justification: Large language models were used to assist with code writing (analysis scripts)
1328 and manuscript editing (grammar, formatting). LLMs were *not* used as a component of
1329 the research methodology itself—all experimental results are derived from deterministic
1330 forward passes through the studied models.